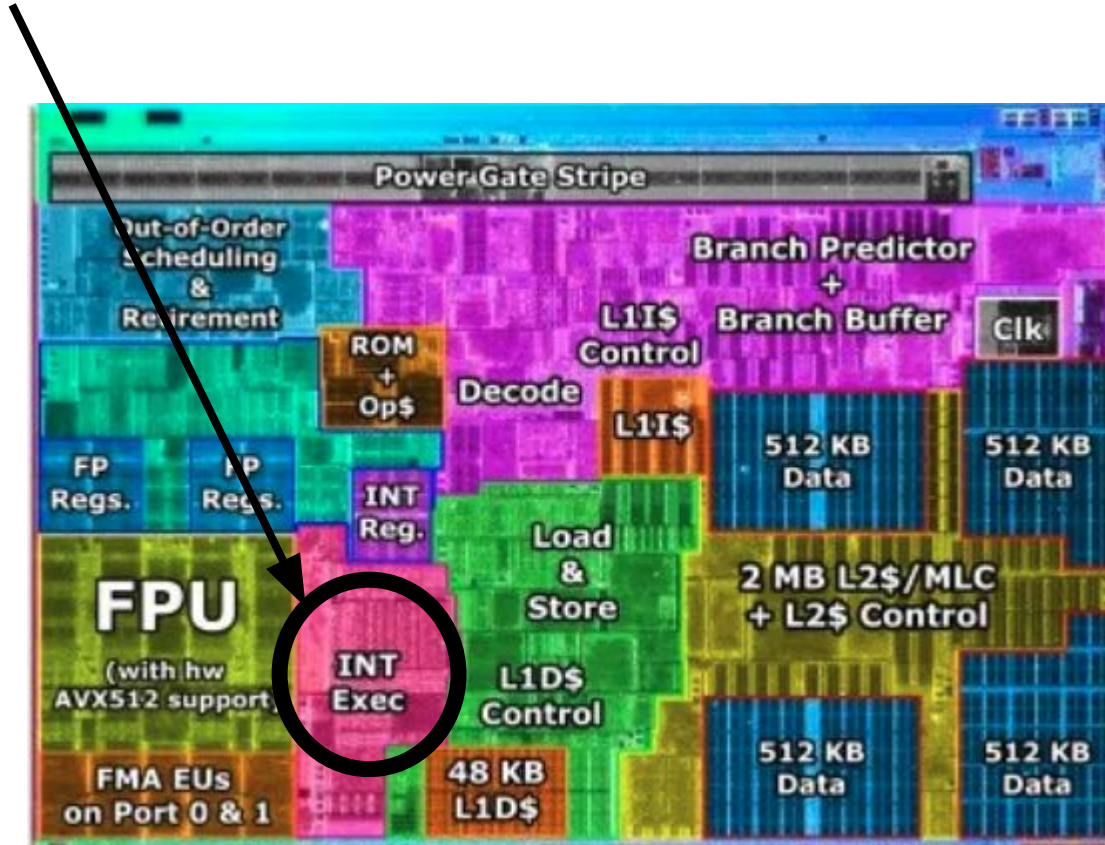


Interface Internals

Keith Randall
@GopherCon, 2024/07/10

```
c := a * b
```

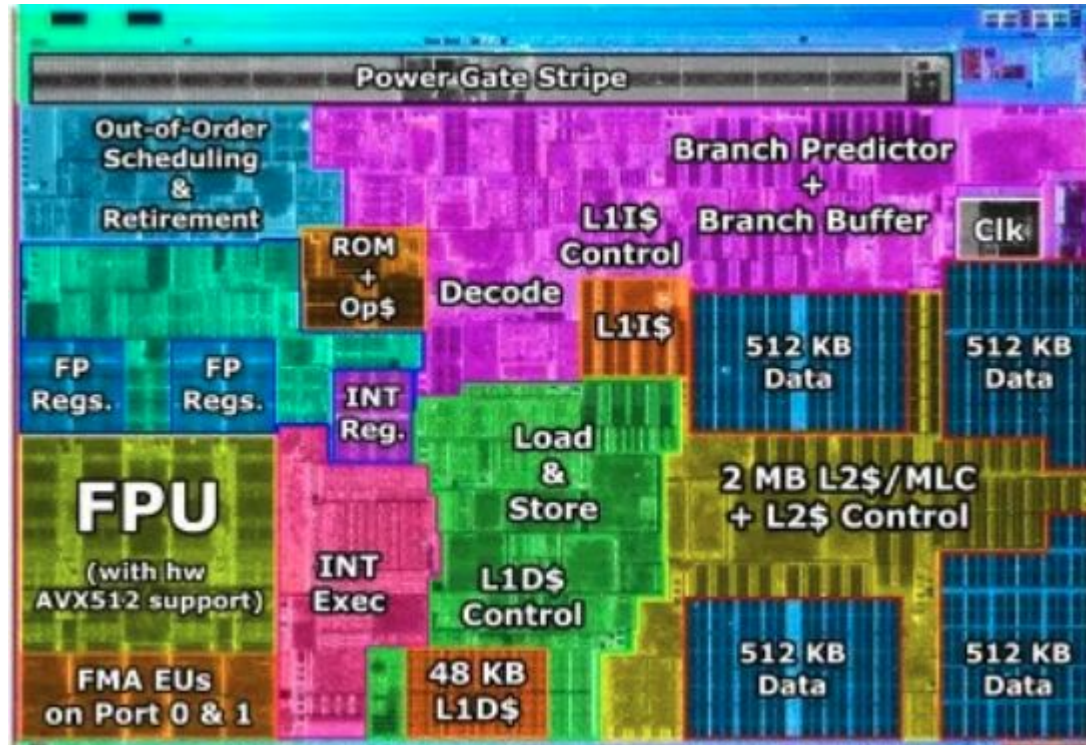
in here somewhere



`c := a * b` \longrightarrow `MUL R1, R0, R0`

```
var c io.Closer = ...  
err := c.Close()
```

There is no “interface method execution unit” anywhere here.



```
var c io.Closer = ...  
err := c.Close()
```

```
var c io.Closer = ...  
err := c.Close()
```

- find the contained type
- get its list of methods
- find the one with the right name
- get its location in the binary
- jump to it


```
var c io.Closer = ...  
err := c.Close()
```

- find the contained type
- get its list of methods
- find the one with the right name
- get its location in the binary
- jump to it

How many instructions will it take?

```
var c io.Closer = ...  
err := c.Close()
```



```
MOVD    24(R0), R2  
MOVD    R1, R0  
CALL    (R2)
```

interface { Read(buf []byte) (int, error) }	io.Reader
interface { Read(buf []byte) (int, error) Write(buf []byte) (int, error) }	io.ReadWriter
interface { Error() string }	a.k.a. error
interface { }	a.k.a. any

Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...}
```

Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...}
```

sure!

Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...}           sure!
```

```
r = &bytes.Buffer{...}
```

Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...}           sure!
```

```
r = &bytes.Buffer{...}      yep!
```

Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...}           sure!
```

```
r = &bytes.Buffer{...}      yep!
```

```
r = 9
```



Interface-typed variables can contain any type that has the methods listed in the interface.

```
var r io.Reader
```

```
r = &os.File{...} // sure!
```

```
r = &bytes.Buffer{} // yep!
```

```
r = 9
```



Interface variables bridge the gap between static and dynamic worlds.

Static

- compile time
- fixed type
- interface type

Dynamic

- run time
- type can change
- value can change
- non-interface (“concrete”) type

Interface variables bridge the gap between static and dynamic worlds.

Static

- compile time
- fixed type
- interface type

Dynamic

- run time
- **type can change**
- value can change
- non-interface (“concrete”) type




This is what distinguishes interfaces
from all other types in Go.

```
var r io.Reader
r = &os.File{...}
n, err := r.Read(buf)
r = &bytes.Buffer{...}
n, err = r.Read(buf)
```

```
var r io.Reader
r = &os.File{...}
n, err := r.Read(buf)
r = &bytes.Buffer{...}
n, err = r.Read(buf)
```

calls os.(*File).Read



```
var r io.Reader
r = &os.File{...}
n, err := r.Read(buf)
r = &bytes.Buffer{...}
n, err = r.Read(buf)
```

calls os.(*File).Read



calls bytes.(*Buffer).Read



This all sounds kind of magical!



Let's investigate how it is done!

Conceptually, interfaces contain a pair of

- a concrete (non-interface) type
- a value of that type

How do we represent that at runtime?

Interfaces are just 2-word structs!

From the runtime:

```
type eface struct {  
    _type *_type  
    data  unsafe.Pointer  
}
```

Interfaces are just 2-word structs!

From the runtime:

```
type eface struct {  
    _type *_type  
    data  unsafe.Pointer  
}
```

↖ We'll tackle the data field first.

```
var i any
```


```
i = [3]int{1, 2, 3}
```


```
var i any
i = [3]int{1, 2, 3}
```

How do we fit a value that big into a 2-word interface?


```
var i any
i = [3]int{1, 2, 3}
```

 = registers / stack


 = heap



```
var i any
i = [3]int{1, 2, 3}
i = [4]float64{4.0, 5.0, 6.0, 7.0}
```



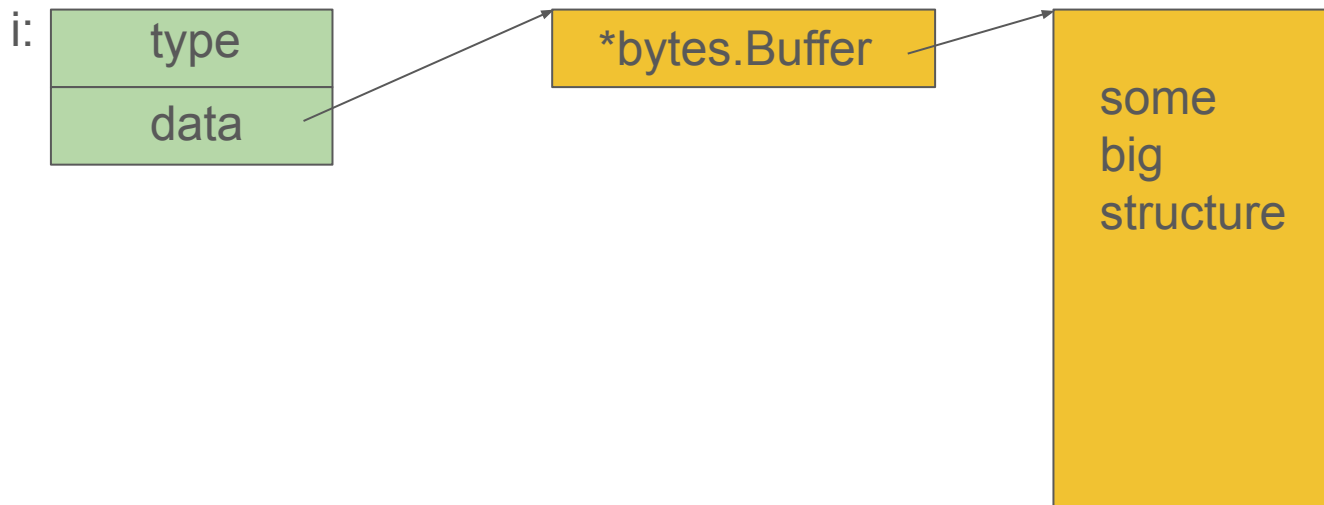
```
var i any
i = [3]int{1, 2, 3}
i = [4]float64{4.0, 5.0, 6.0, 7.0}
```



Optimization: if the data is already a pointer, we can use it directly.

```
var i any
i = &bytes.Buffer{...}
```


instead of



Optimization: if the data is already a pointer, we can use it directly.

```
var i any  
i = &bytes.Buffer{...}
```


we can do



Optimization: if the data is already a pointer, we can use it directly.

```
var i any
i = &bytes.Buffer{...}
```

we can do



allocation-free assigning to interfaces,
if the concrete type is a pointer!

Interfaces are just 2-word structs!


From the runtime:

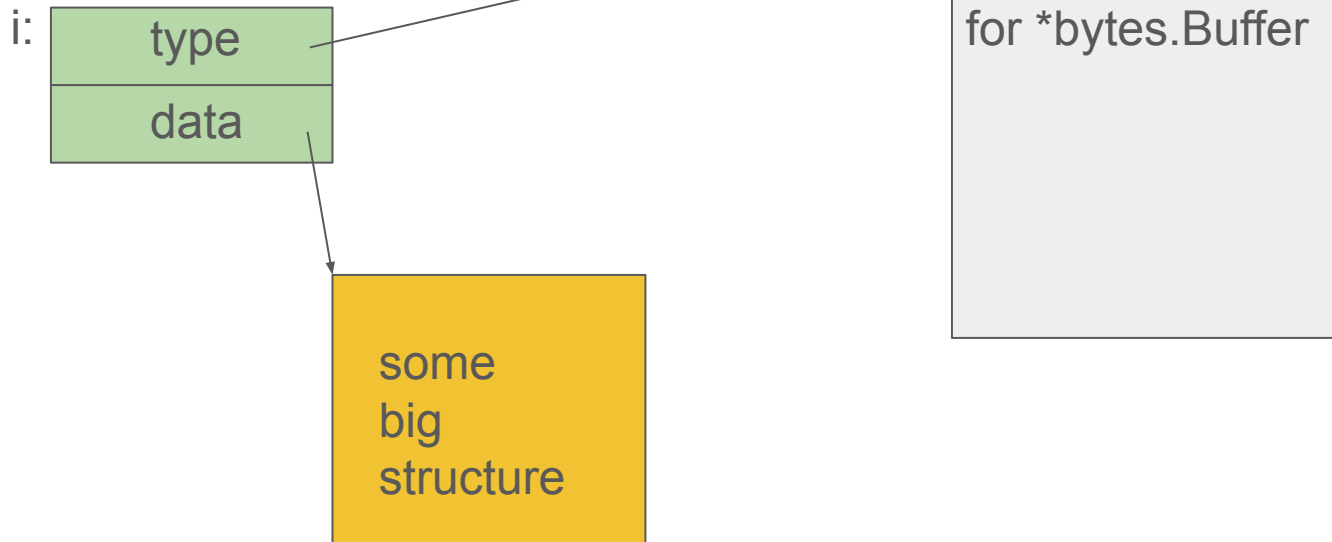
```
type eface struct {  
    _type *_type  
    data  unsafe.Pointer  
}
```

Now we'll look at the type information.



```
var i any
i = &bytes.Buffer{...}
```

 = static



Type descriptors

Give the runtime access to information about a type.

- Size
- Pointer fields
- String representation
- How to do `==` on this type
- List of methods
- Lives in the read-only data section of the binary
 - or maybe allocated by the reflect package

concrete -> interface conversion

```
var b *bytes.Buffer = ...  
var i any = b
```

... b starts in R0 ...

```
MOVD    R0, R1                // i.data = b  
MOVD    $type:*bytes.Buffer(SB), R0  // i._type = constant  
... i is in R0,R1 ...
```

interface -> concrete conversion

```
var i any = ...  
if b, ok := i.(*bytes.Buffer); ok { ... }
```

... i starts in R0,R1 ...

```
MOVD    $type:*bytes.Buffer(SB), R2    // R2 = constant  
CMP     R2, R0                          // R2 == i._type?  
BEQ     32                               // branch if equal  
... b is in R1 ...
```

interface -> concrete conversion

- Type descriptors must be unique.
- Deduplicated by the linker.
- The **reflect** and **plugin** packages have to be careful not to break this property.

Compare against nil

```
err := f()  
if err != nil { ... }
```

... err is in R0/R1 ...

```
CBNZ    R0, 48          // branch if err._type != nil
```

```
var c io.Closer = ...  
err := c.Close()
```

```
var c io.Closer = ...  
err := c.Close()
```

- find the contained type
- get its list of methods
- find the one with the right name
- get its location in the binary
- jump to it


```
var c io.Closer = ...  
err := c.Close()
```

- ~~find the contained type~~ done!
- get its list of methods
- find the one with the right name
- get its location in the binary
- jump to it

```
var c io.Closer = ...  
err := c.Close()
```

- ~~find the contained type~~ done!
- get its list of methods
- find the one with the right name
- get its location in the binary
- jump to it

The rest of this work can be expensive!
Can we precompute it somehow?




Interfaces are just 2-word structs!

From the runtime:


```
type eface struct {  
    _type *_type  
    data  unsafe.Pointer  
}
```

Problem: there's no room!

i:



type descriptor
for *bytes.Buffer



- The intermediate object is called an “interface table”, or “itab”.
- Provides space to store additional data.

Interfaces are just 2-word structs!

From the runtime:

```
type eface struct  
    _type *_type  
    data  unsafe.Pointer  
}
```



Interfaces are just 2-word structs!

Empty interfaces, from the runtime:

```
type eface struct {  
    _type *_type  
    data  unsafe.Pointer  
}
```

Non-empty interfaces, from the runtime:

```
type iface struct {  
    tab *_tab  
    data unsafe.Pointer  
}
```

Interface Tables

```
type itab struct {  
    _type *_type  
    inter *interfaceType  
    fun    [1]uintptr // variable sized.  
}
```

Interface Tables

```
type itab struct {  
    _type *_type ← dynamic, concrete type  
    inter *interfaceType  
    fun    [1]uintptr // variable sized.  
}
```

Interface Tables

```
type itab struct {  
    _type *_type  
    inter *interfaceType ← static, interface type  
    fun    [1]uintptr // variable sized.  
}
```

Interface Tables

```
type itab struct {  
    _type *_type  
    inter *interfaceType  
    fun    [1]uintptr // variable sized.  
}
```



“method table”, 1 entry for each method of
the interface (in sorted order)

Interface Tables

```
type itab struct {  
    _type *_type  
    inter *interfaceType  
    fun    [1]uintptr // variable sized.  
}
```

each entry is a PC for the start of a method



Interface Tables

```
var r io.Reader = &bytes.Buffer{}
```


itab:*bytes.Buffer,io.Reader

_type
inter
fun[0]

type:*bytes.Buffer



type:io.Reader



bytes.(*Buffer).Read

```
MOVD 16(R28), R16
CMP  R16, RSP
BLS  65(PC)
MOVD.W R30, -48(RSP)
MOVD  R29, -8(RSP)
SUB   $8, RSP, R29
MOVD  R1, 64(RSP)
MOVB  ZR, 32(R0)
...
```


Interface Tables

```
var r io.Reader = &bytes.Buffer{}
```

itab:*bytes.Buffer,io.Reader


_type
inter
fun[0]

All static data in the binary, all built by the compiler.

type:*bytes.Buffer



type:io.Reader



bytes.(*Buffer).Read

```
MOVD 16(R28), R16
CMP  R16, RSP
BLS  65(PC)
MOVD.W R30, -48(RSP)
MOVD  R29, -8(RSP)
SUB   $8, RSP, R29
MOVD  R1, 64(RSP)
MOVB  ZR, 32(R0)
...
```


Interface Tables

```
var r io.ReadWriter = &bytes.Buffer{}
```

itab:*bytes.Buffer,io.ReadWriter

_type
inter
fun[0]
fun[1]

type:*bytes.Buffer



type:io.ReadWriter



bytes.(*Buffer).Write

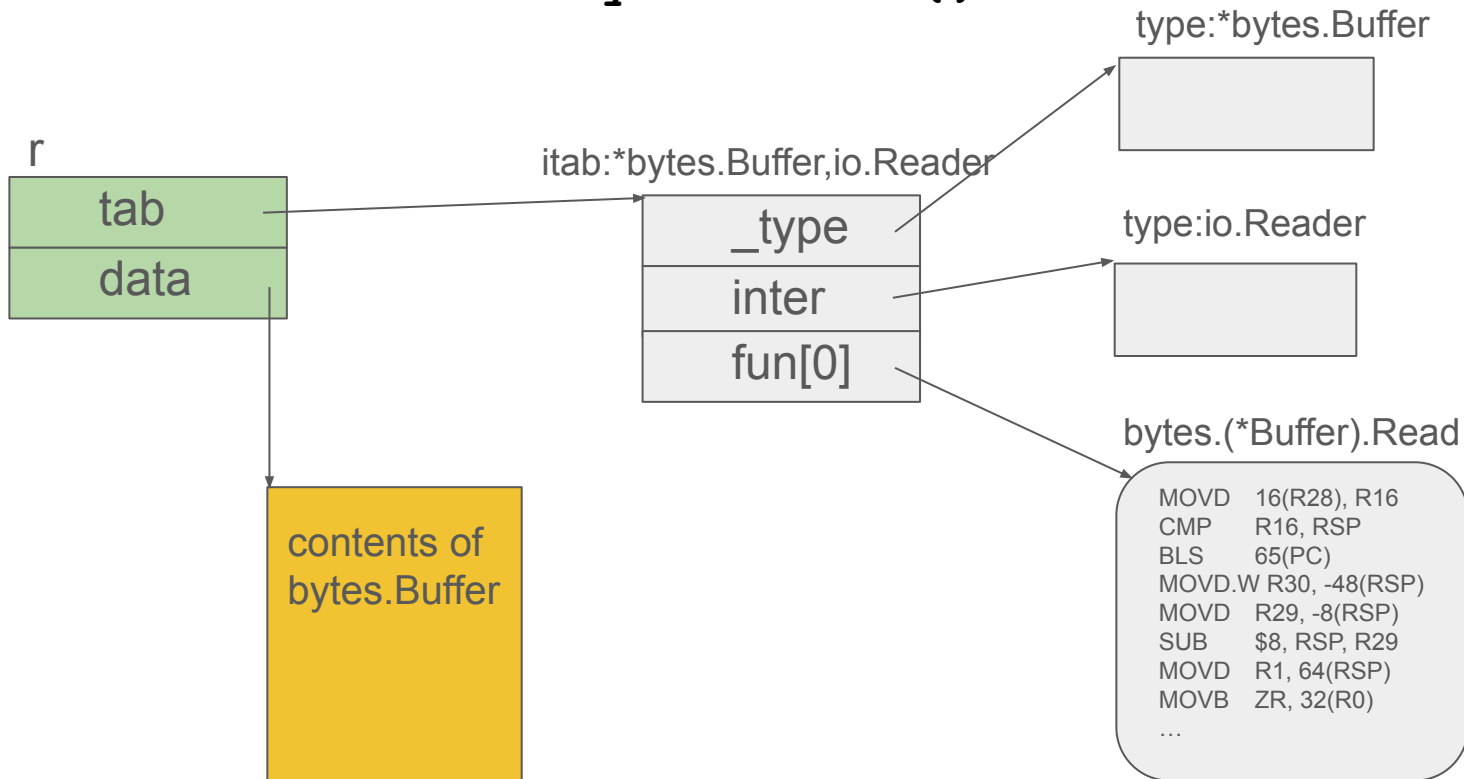
```
MOVD 16(R28), R16
CMP  R16, RSP
BLS  55(PC)
MOVD.W R30, -48(RSP)
MOVD  R29, -8(RSP)
SUB   $8, RSP, R29
MOVD  R1, 64(RSP)
MOVB  ZR, 32(R0)
...
```

bytes.(*Buffer).Read

```
MOVD 16(R28), R16
CMP  R16, RSP
BLS  65(PC)
MOVD.W R30, -48(RSP)
MOVD  R29, -8(RSP)
SUB   $8, RSP, R29
MOVD  R1, 64(RSP)
MOVB  ZR, 32(R0)
...
```

Interface Tables

```
var r io.Reader = &bytes.Buffer{}
```

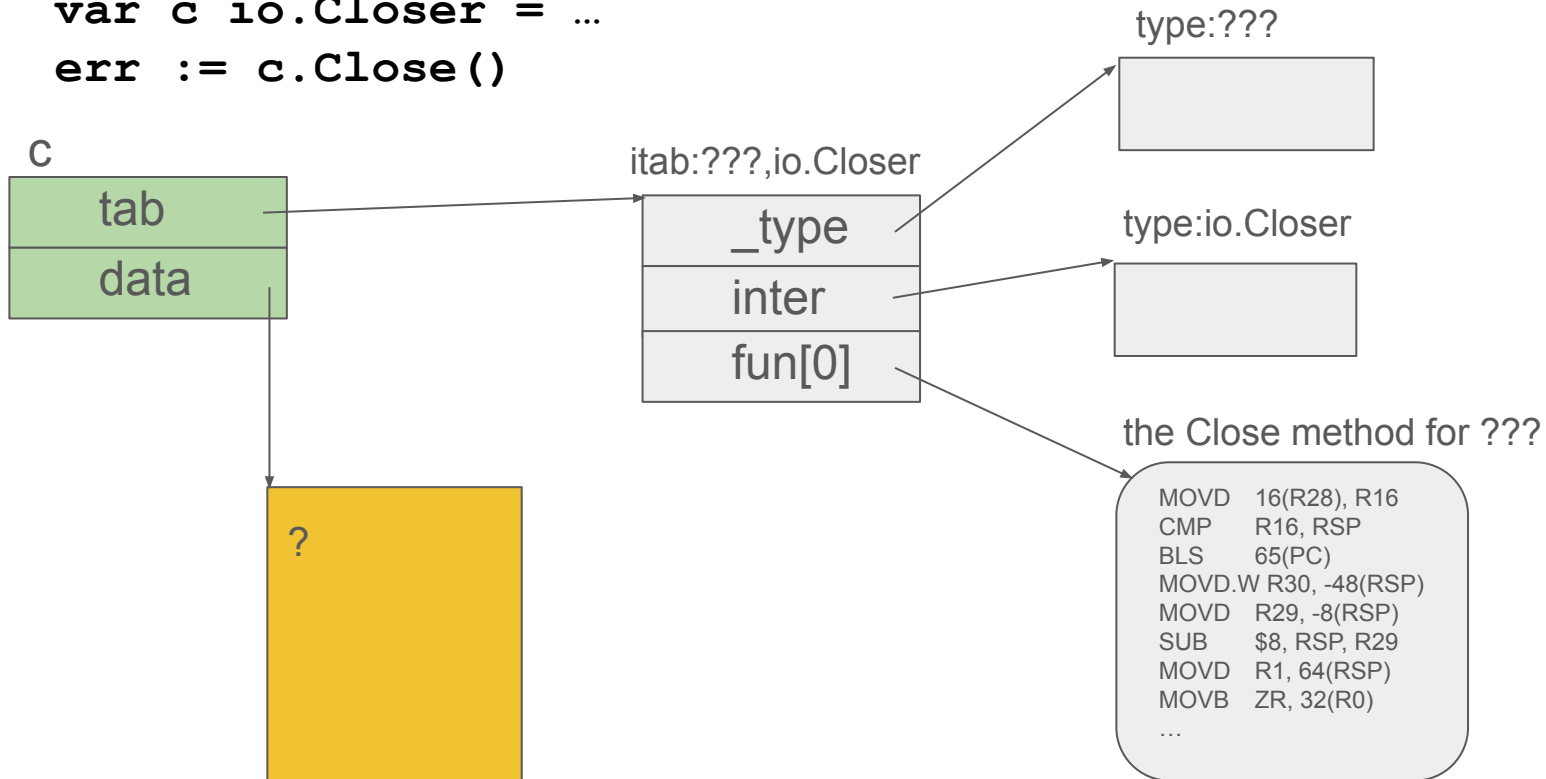


interface method invocation

```
var c io.Closer = ...  
err := c.Close()
```

interface method invocation

```
var c io.Closer = ...  
err := c.Close()
```



interface method invocation

```
var c io.Closer = ...  
err := c.Close()
```

... c is in R0,R1 ...

```
MOVD    24(R0), R2
```

```
MOVD    R1, R0
```

```
CALL    (R2)
```

... Close method's receiver is in R0 ...

```
// pc := c.tab.fun[0]
```

```
// receiver := c.data
```

```
// jump to pc
```

interface method invocation

```
var c io.Closer = ...  
err := c.Close()
```

24 = offset of fun[0] in the itab struct

... c is in R0,R1 ...

```
MOVD    24(R0), R2
```

```
MOVD    R1, R0
```

```
CALL    (R2)
```

... Close method's receiver is in R0 ...

// pc := c.tab.fun[0]

// receiver := c.data

// jump to pc

interface -> interface conversion

go1.21 and earlier

```
var r io.Reader = ...  
if rw, ok := r.(io.ReadWriter); ok { ... }
```

... r is in R0,R1 ...

```
MOVD    R1, R2  
MOVD    R0, R1  
MOVD    $type:io.ReadWriter(FP), R0  
CALL    runtime.assertI2I2(SB)  
CBZ     R0, 64
```

... rw is in R0/R1 ...

interface -> interface conversion

go1.21 and earlier

```
var r io.Reader = ...  
if rw, ok := r.(io.ReadWriter); ok { ... }
```

... r is in R0,R1 ...

MOVD R1, R2

MOVD R0, R1

MOVD \$type:io.ReadWriter(FP), R0

CALL runtime.assertI2I2(SB)

CBZ R0, 64

... rw is in R0/R1 ...

The tricky part: we need
a new interface table for
the result



interface -> interface conversion

go1.21 and earlier

The runtime call:

1. Checks a cache to see if we've made this interface table before.
 - a. *Required* because interface tables must be unique.
2. Obtains a list of all the methods of the concrete type.
3. Finds each of the interface's methods in that list.
4. Builds an interface table, caches it, and returns it.

interface -> interface conversion

go1.22 and later

```
var r io.Reader = ...  
if rw, ok := r.(io.ReadWriter); ok { ... }
```

... r is in R0,R1 ...

... rw is in R0/R1 ...

```
CBZ    R0, 64  
MOVD   8(R0), R2  
MOVD   $main..typeAssert.0(SB), R3  
LDAR   (R3), R4  
MOVWU  16(R0), R5  
MOVD   (R4), R6  
AND    R6, R5, R7  
LSL    $4, R7, R7  
ADD    $8, R7, R7  
MOVD   (R4)(R7), R8  
ADD    R7, R4, R7  
CMP    R2, R8  
BEQ    136  
ADD    $1, R5, R5  
CBNZ   R8, 76  
MOVD   R1, main.rc+8(FP)  
MOVD   R3, R0  
MOVD   R2, R1  
CALL   runtime.typeAssert(SB)  
MOVD   main.rc+8(FP), R1  
JMP    64  
MOVD   8(R7), R0
```

interface -> interface conversion

go1.22 and later

```
var r io.Reader = ...  
if rw, ok := r.(io.ReadWriter); ok { ... }
```

... r is in R0,R1 ...

... rw is in R0/R1 ...

Per call site cache

```
CBZ    R0, 64  
MOVD   8(R0), R2  
MOVD   $main..typeAssert.0(SB), R3  
LDAR   (R3), R4  
MOVWU  16(R0), R5  
MOVD   (R4), R6  
AND    R6, R5, R7  
LSL    $4, R7, R7  
ADD    $8, R7, R7  
MOVD   (R4)(R7), R8  
ADD    R7, R4, R7  
CMP    R2, R8  
BEQ    136  
ADD    $1, R5, R5  
CBNZ   R8, 76  
MOVD   R1, main.rc+8(FP)  
MOVD   R3, R0  
MOVD   R2, R1  
CALL   runtime.typeAssert(SB)  
MOVD   main.rc+8(FP), R1  
JMP    64  
MOVD   8(R7), R0
```

interface -> interface conversion

go1.22 and later

```
var r io.Reader = ...  
if rw, ok := r.(io.ReadWriter); ok { ... }
```

name	old time/op	new time/op	delta
TypeAssert	3.78ns \pm 3%	1.00ns \pm 1%	-73.53%
TypeSwitch	25.8ns \pm 2%	2.5ns \pm 3%	-90.43%


Conclusion

- Interfaces are magic!




Conclusion

- Interfaces are mag



Conclusion

- Interfaces are good engineering.



Things to cover

Interface equality

type asserts, type switches

nil interface vs nil pointer

new faster type switches?

allocation requirements (use pointer types when you can)

interfaces are value types - they can't be changed except wholesale

method invocation (+ itabs)

type descriptor

Skipping

empty \leftrightarrow nonempty conversions

comparing 2 interfaces to each other?

wrapper functions for non-pointers

concrete -> interface conversion

```
var b *bytes.Buffer = ...  
var r io.Reader = b
```

... b starts in R0 ...

```
MOVD    R0, R1
```

```
MOVD    $itab:*bytes.Buffer,io.Reader(SB), R0
```

... r is in R0,R1 ...

interface -> concrete conversion

```
var r io.Reader = ...  
if b, ok := r.(*bytes.Buffer); ok { ... }
```

... r is in R0,R1 ...

```
MOVD    $itab:*bytes.Buffer,io.Reader(SB), R2  
CMP     R2, R0  
BNE     32  
... b is in R1 ...
```

interface -> any conversion

```
var r io.Reader = ...  
var i any = r
```

... r is in R0,R1 ...

```
CBZ    R0, 16
```

```
MOVD   8(R0), R0
```

... i is in R0/R1 ...

interface -> any conversion

```
var r io.Reader = ...  
var i any = r
```

... r is in R0,R1 ...

```
CBZ    R0, 16
```

```
MOVD   8(R0), R0
```

... i is in R0/R1 ...

← copies the Type field out of
the interface table

data field unmodified!

interface -> interface down conversion

go1.21 and earlier

```
var rc io.ReadCloser = ...  
var r io.Reader = rc
```

... rc is in R0,R1 ...

```
MOVD    R1, main.rc+8(FP)  
MOVD    R0, R1  
MOVD    $type:io.Closer(FP), R0  
CALL    runtime.convI2I(SB)  
MOVD    main.rc+8(FP), R1  
... r is in R0/R1 ...
```

interface -> interface down conversion

go1.22 and later

```
var rc io.ReadCloser = ...
```

```
var r io.Reader = rc
```

... rc is in R0,R1 ...

... r is in R0/R1 ...

```
CBZ    R0, 64
MOVD   8(R0), R2
MOVD   $main..typeAssert.0(SB), R3
LDAR   (R3), R4
MOVWU  16(R0), R5
MOVD   (R4), R6
AND    R6, R5, R7
LSL    $4, R7, R7
ADD    $8, R7, R7
MOVD   (R4)(R7), R8
ADD    R7, R4, R7
CMP    R2, R8
BEQ    136
ADD    $1, R5, R5
CBNZ   R8, 76
MOVD   R1, main.rc+8(FP)
MOVD   R3, R0
MOVD   R2, R1
CALL   runtime.typeAssert(SB)
MOVD   main.rc+8(FP), R1
JMP    64
MOVD   8(R7), R0
```