

# Building a high-performance concurrent map in Go

YunHao Zhang

ByteDance programming language team

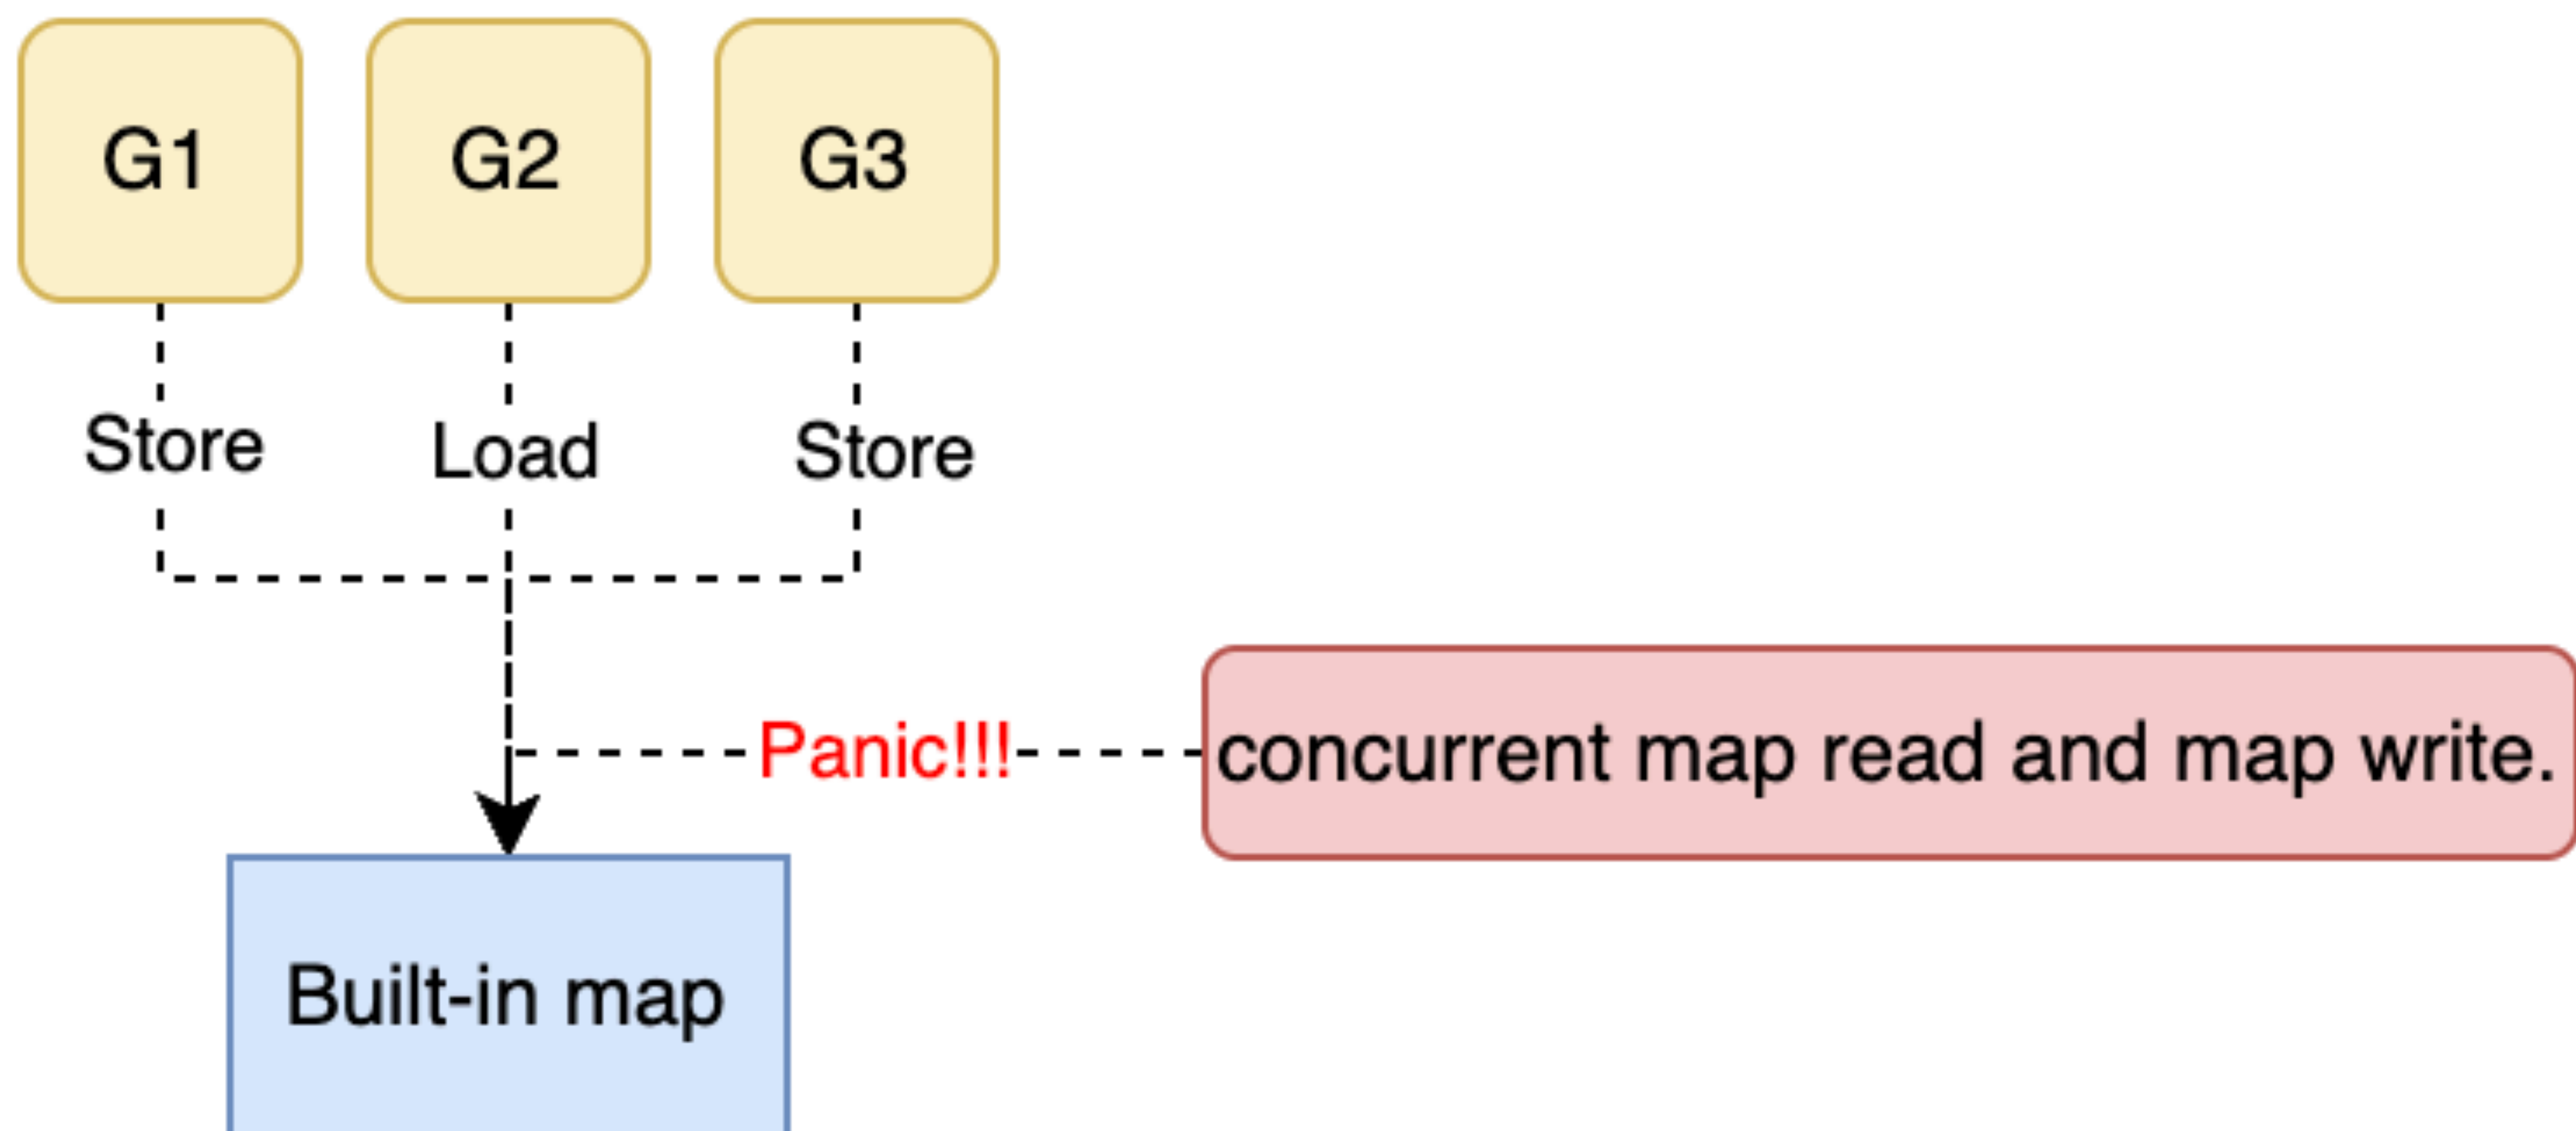
[github.com/zhangyunhao116](https://github.com/zhangyunhao116)



} Key ideas - 25 min

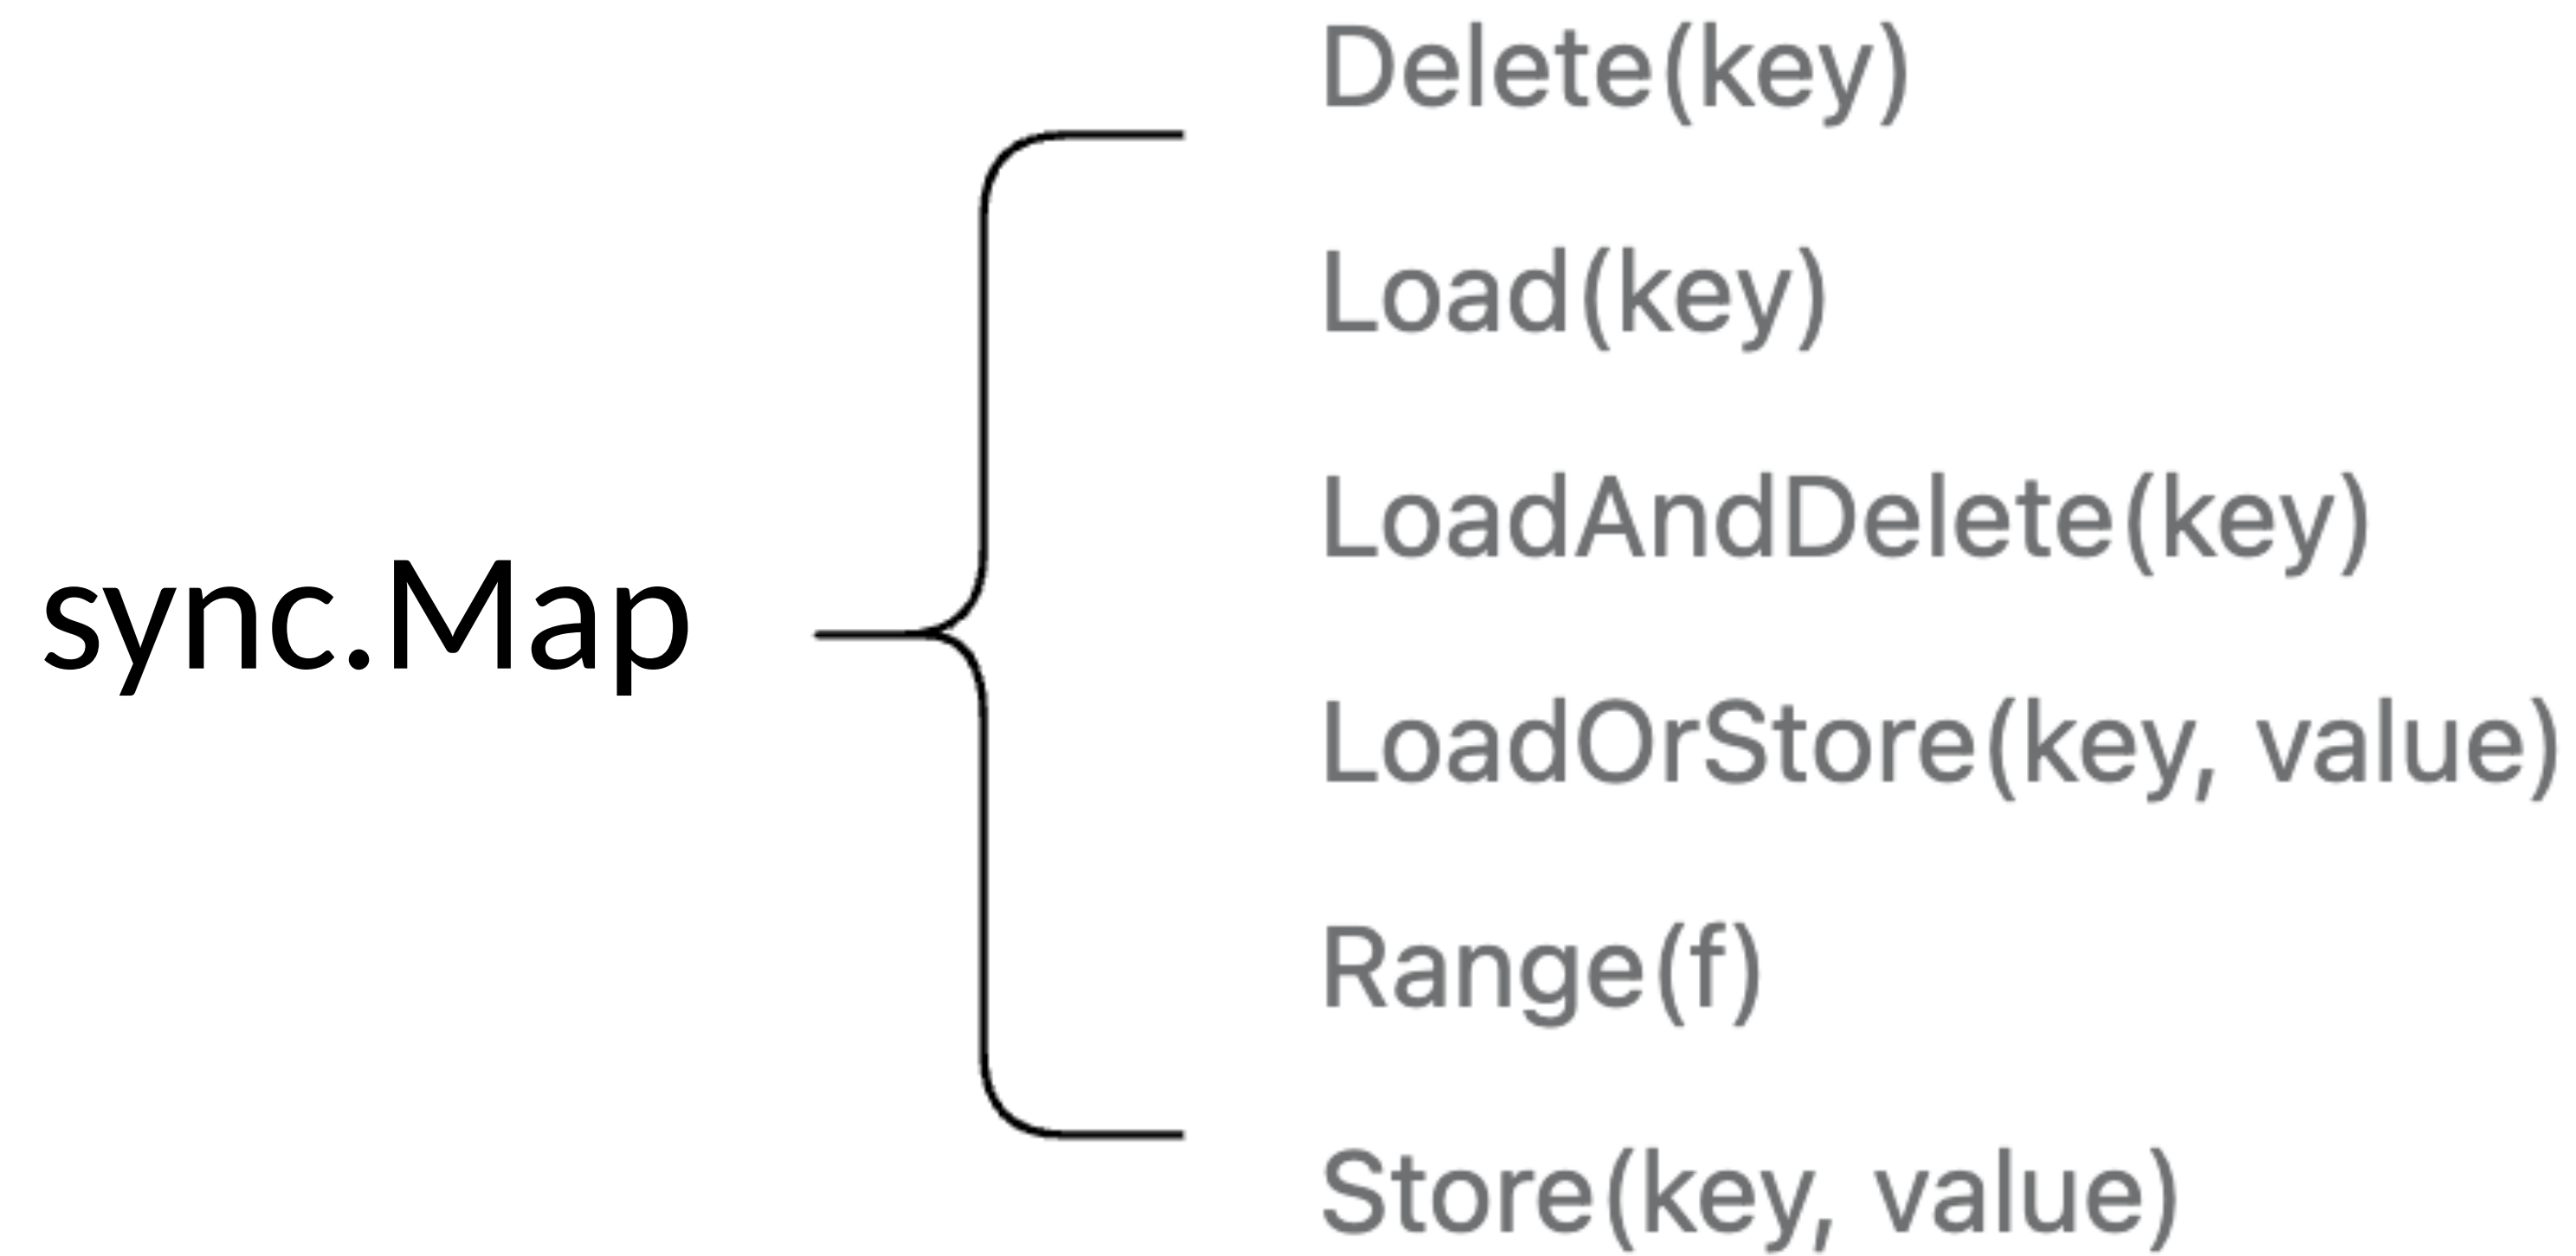
} Details - 2h30 min

Why we need concurrent map ?



The built-in map is NOT concurrent-safe

# How to build a concurrent-map



`map[K]V + sync.RWMutex`

	rwmap.txt sec/op	sec/op	syncmap.txt vs base
LoadSize1000-16	29.565n ± 1%	7.151n ± 23%	-75.81% (p=0.000 n=10)
Store-16	497.0n ± 8%	913.4n ± 3%	+83.77% (p=0.000 n=10)
70Load30Store-16	195.1n ± 5%	699.5n ± 5%	+258.56% (p=0.000 n=10)
90Load9Store1Delete-16	139.7n ± 2%	595.5n ± 1%	+326.42% (p=0.000 n=10)
90Load8Store1Delete1Range-16	8.497μ ± 2%	40.594μ ± 1%	+377.75% (p=0.000 n=10)
geomean	320.9n	643.6n	+100.58%

LoadSize1000: Only Load operation in a map with 1000 items

Store: Only Store operation in a map

70Load30Store: 70% Load and 30% Store operations

90Load9Store1Delete: 90% Load, 9% Store, 1% Delete


90Load8Store1Delete1Range: 90% Load, 8% Store, 1% Delete, 1% Range



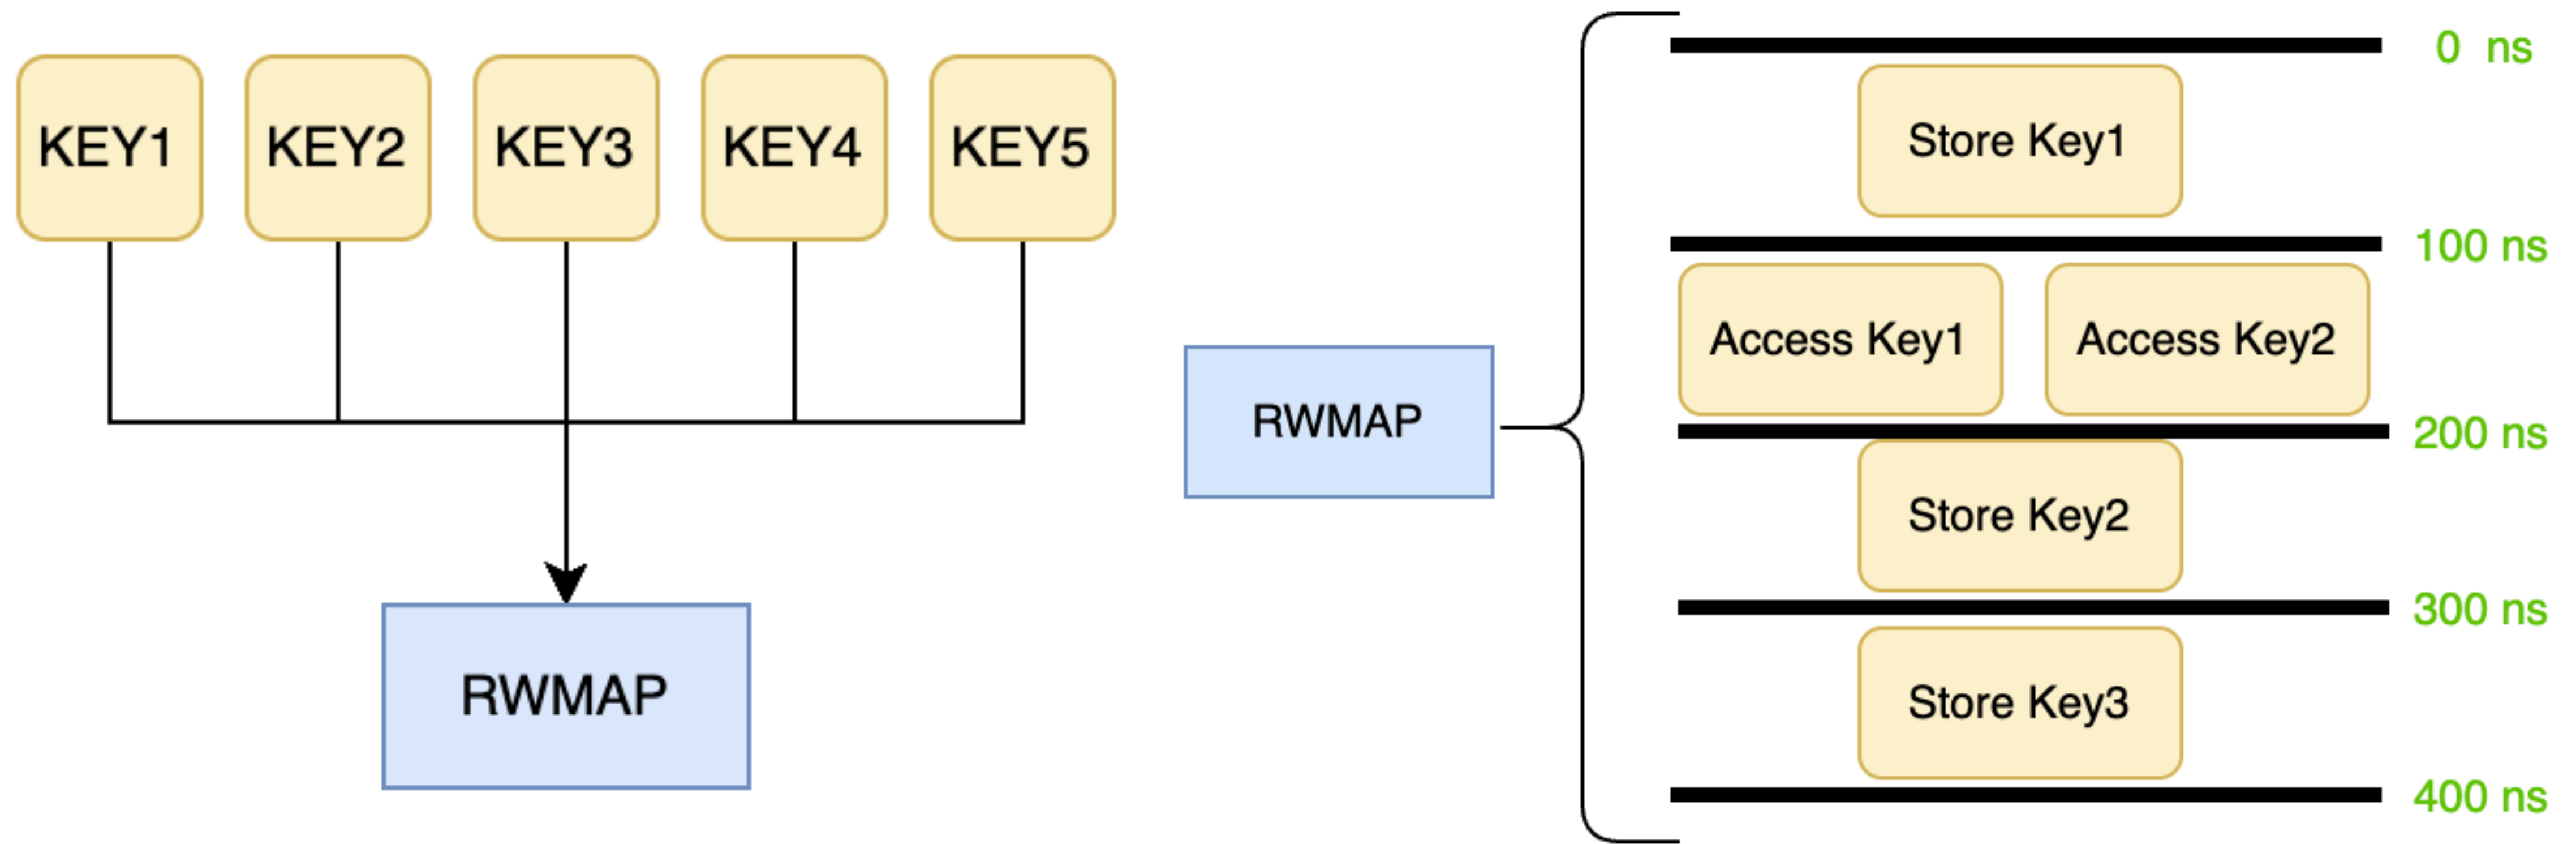
For **Load-only** case, sync.Map is faster

For **Mixed read-write** cases, read-write mutex map is faster

sync.Map is designed to solve the problem known as **cache contention**

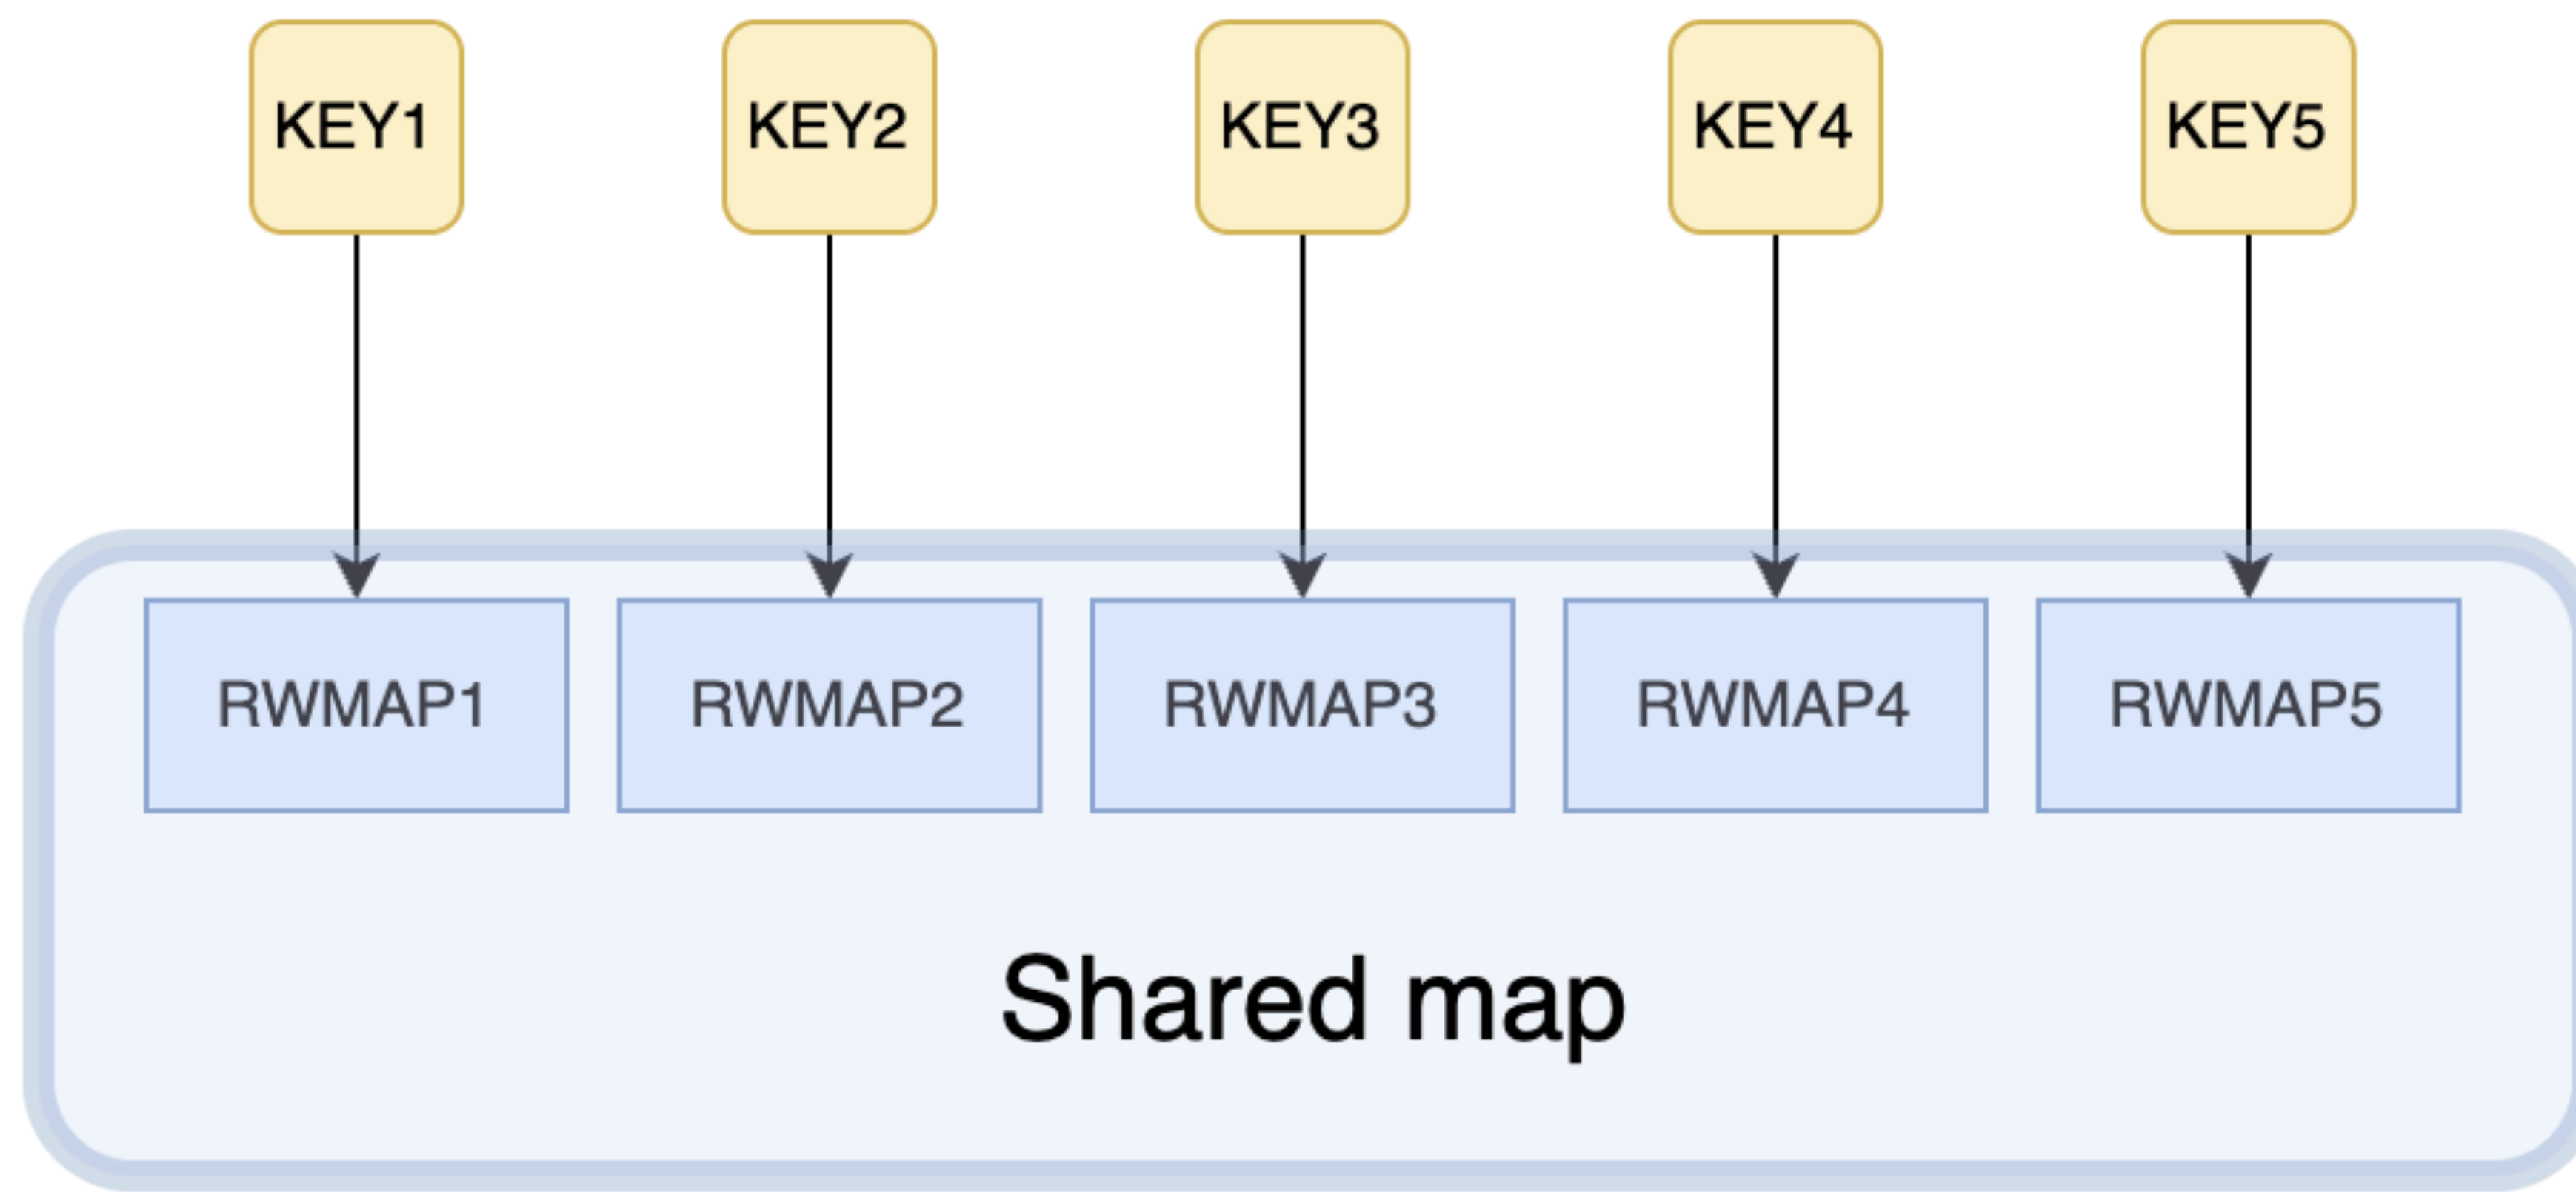


Using the map as a concurrent cache



\*Only one write operation can be executed at a time

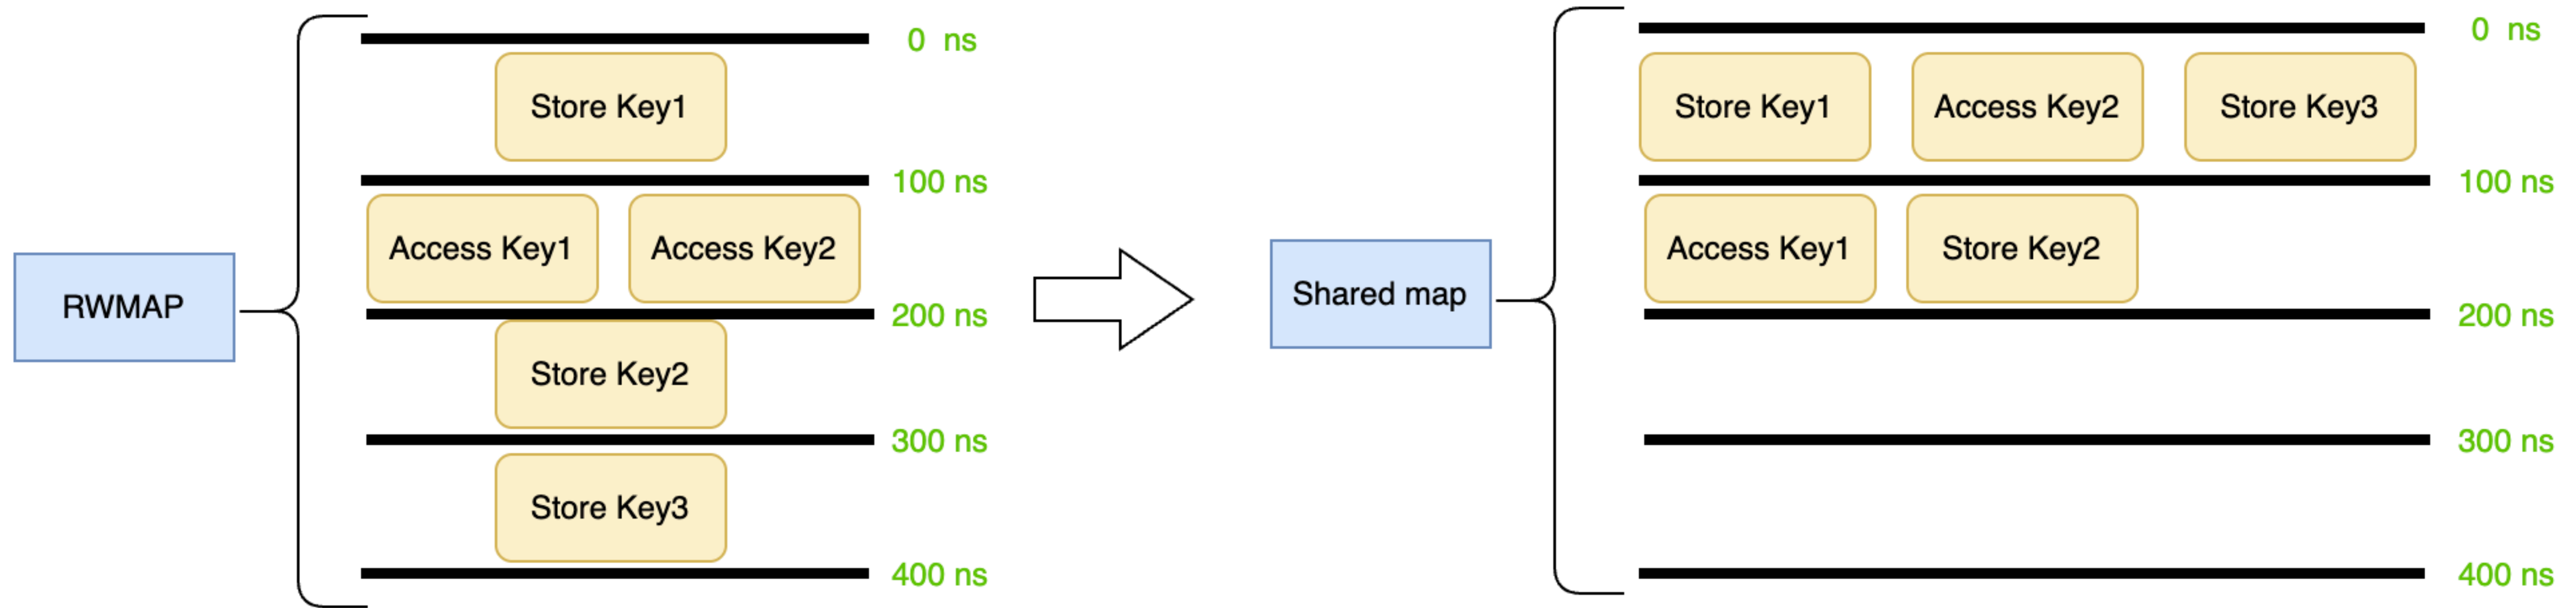
\*Multiple read operations can be executed concurrently



Shared map = multiple read-write mutex maps

\* X CPU cores => X read-write mutex maps in the shared map

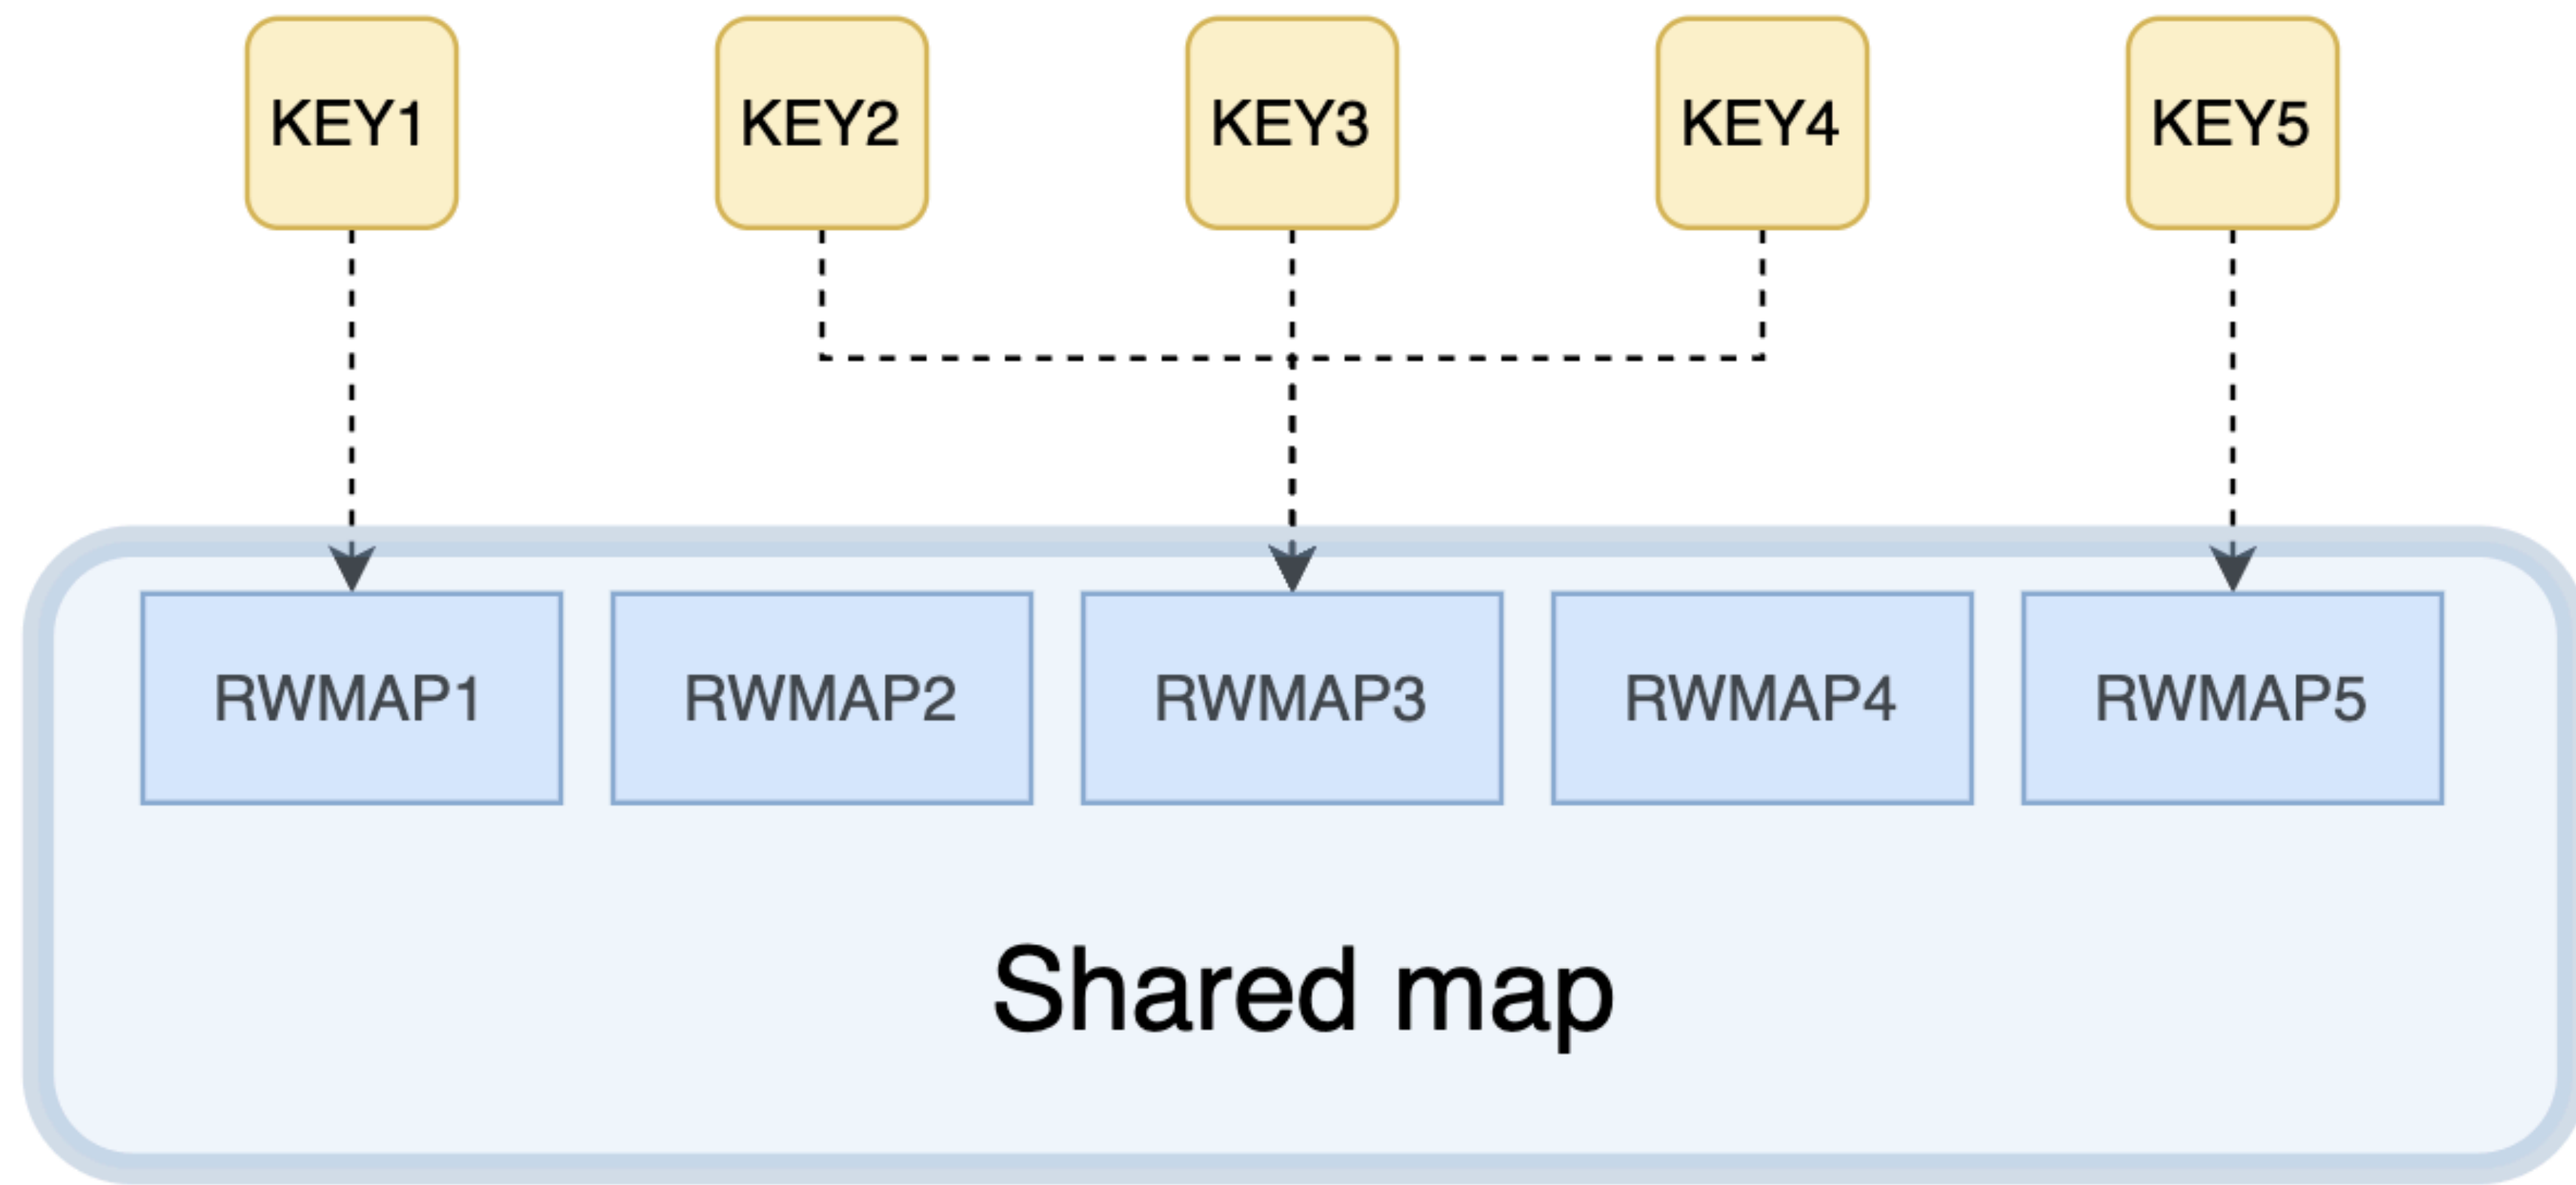
\* Use a hash function to determine which map a key should go to



The single read-write mutex map versus the shared map.

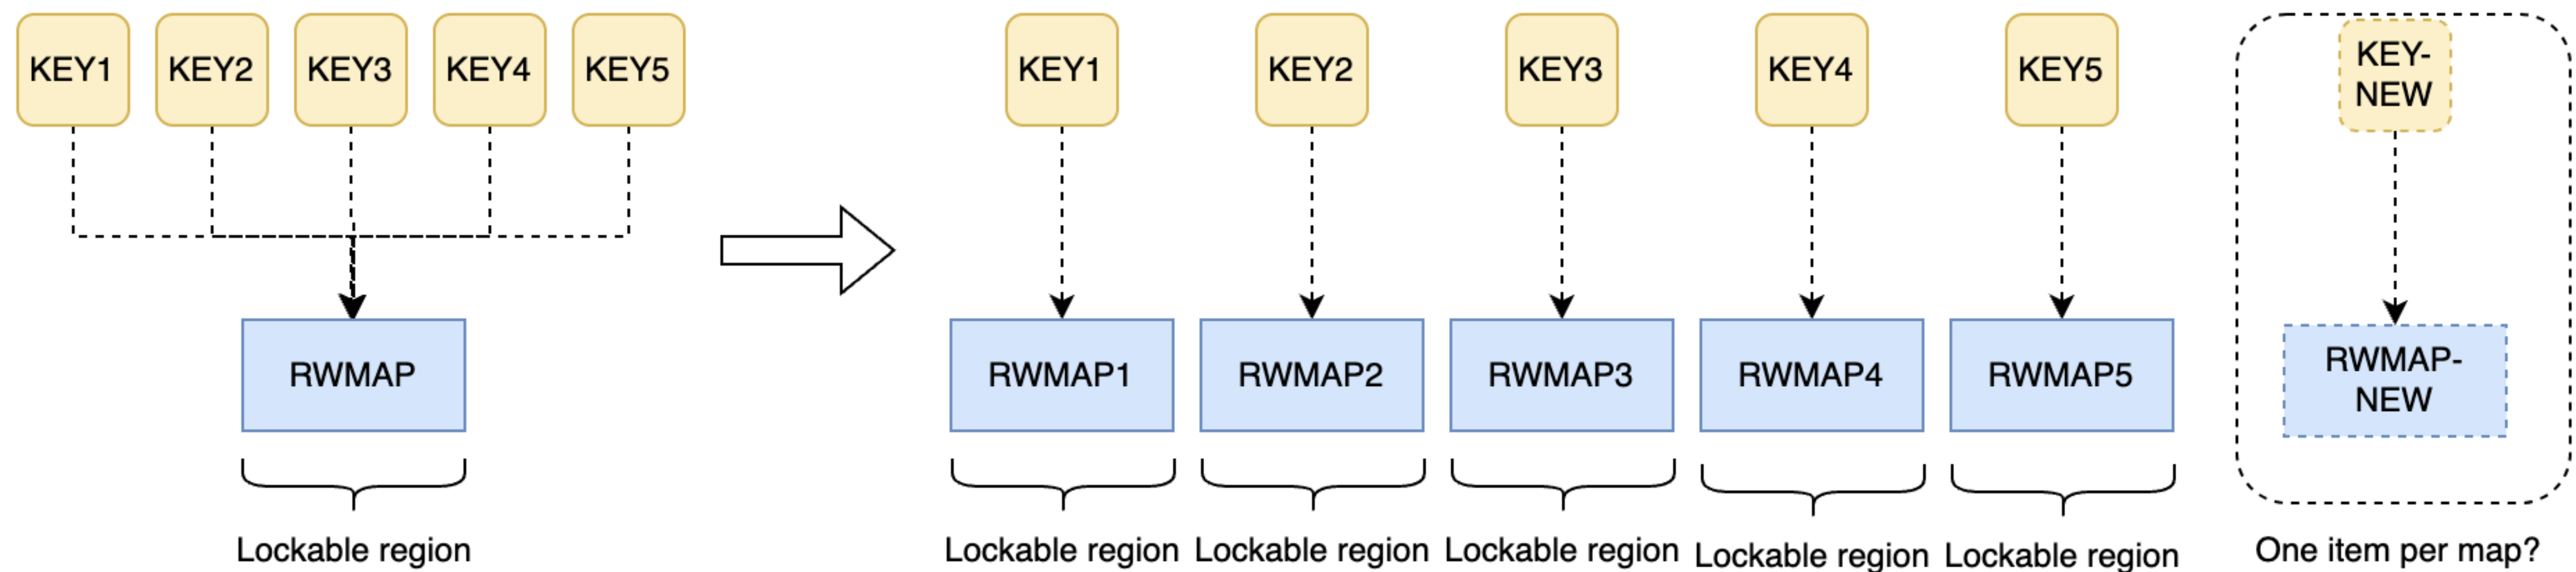
	rwmap.txt	sharedmap.txt		
	sec/op	sec/op	vs base	
LoadSize1000-16	29.56n ± 1%	13.63n ± 11%	-53.88% (p=0.000 n=10)	
Store-16	497.0n ± 8%	130.2n ± 7%	-73.80% (p=0.000 n=10)	
70Load30Store-16	195.1n ± 5%	151.9n ± 2%	-22.17% (p=0.000 n=10)	
90Load9Store1Delete-16	139.65n ± 2%	92.69n ± 3%	-33.63% (p=0.000 n=10)	
90Load8Store1Delete1Range-16	8.497μ ± 2%	1.652μ ± 2%	-80.56% (p=0.000 n=10)	
geomean	320.9n	132.8n	-58.62%	

Shared map is faster in all cases, but it requires more memory



A perfect Hash function doesn't exist!  
(Different operations may still go to the same sub-map)

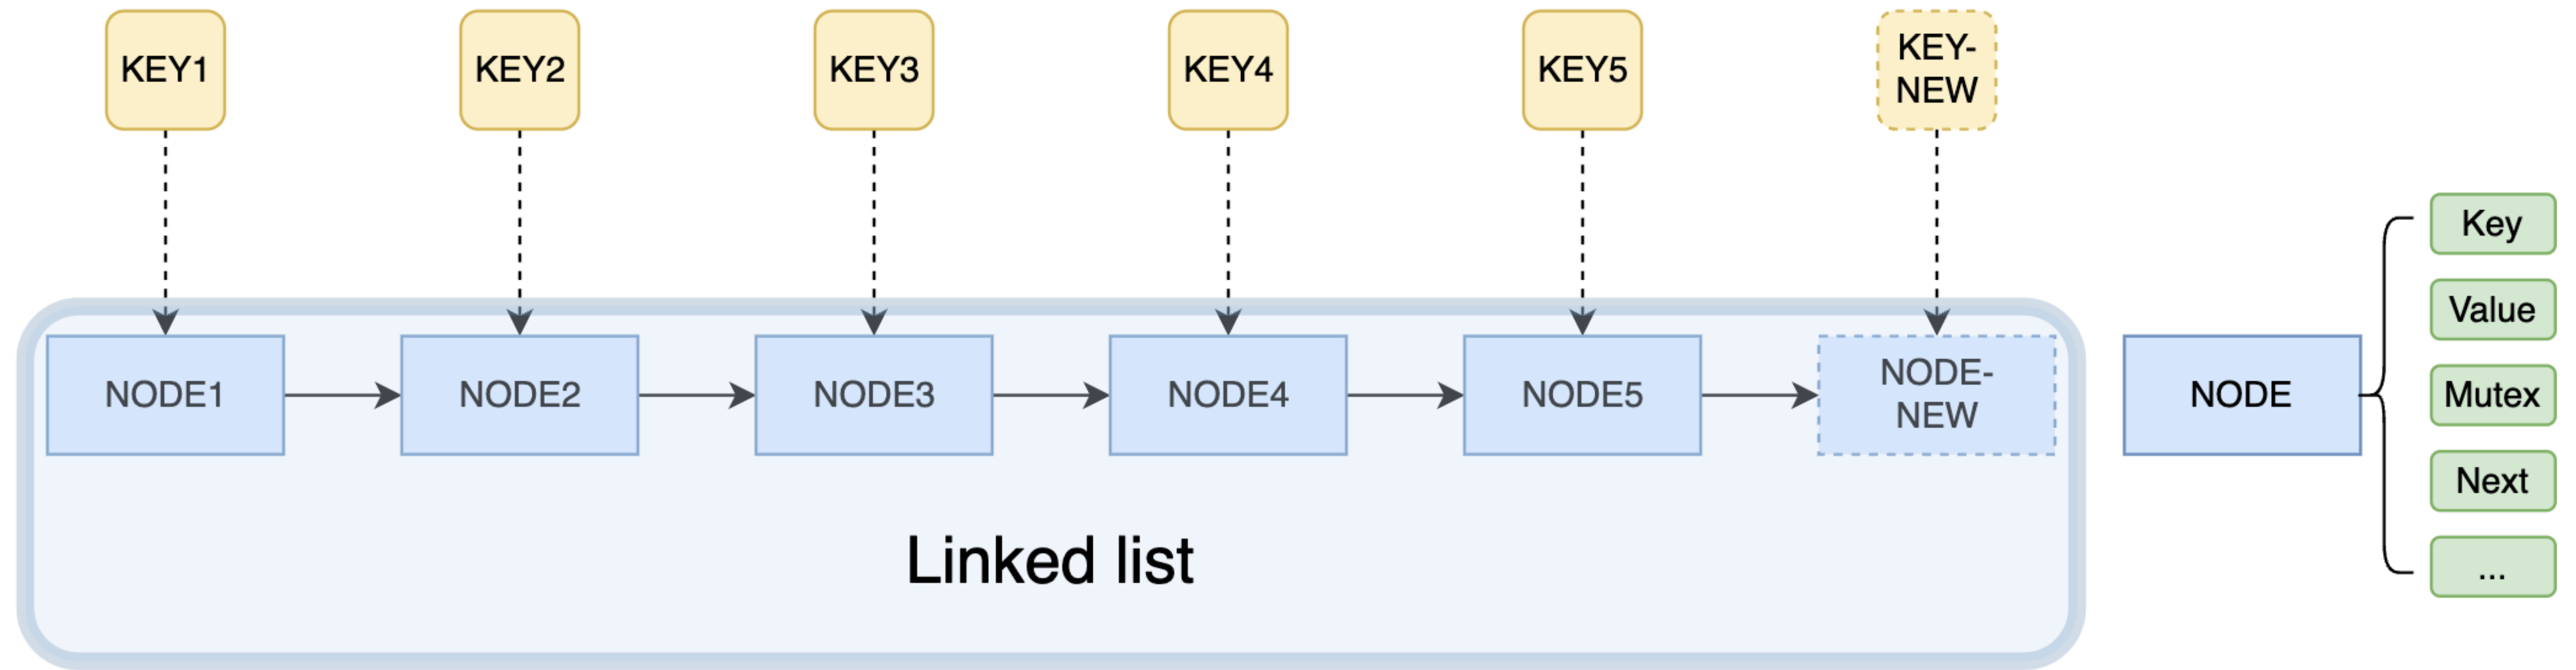




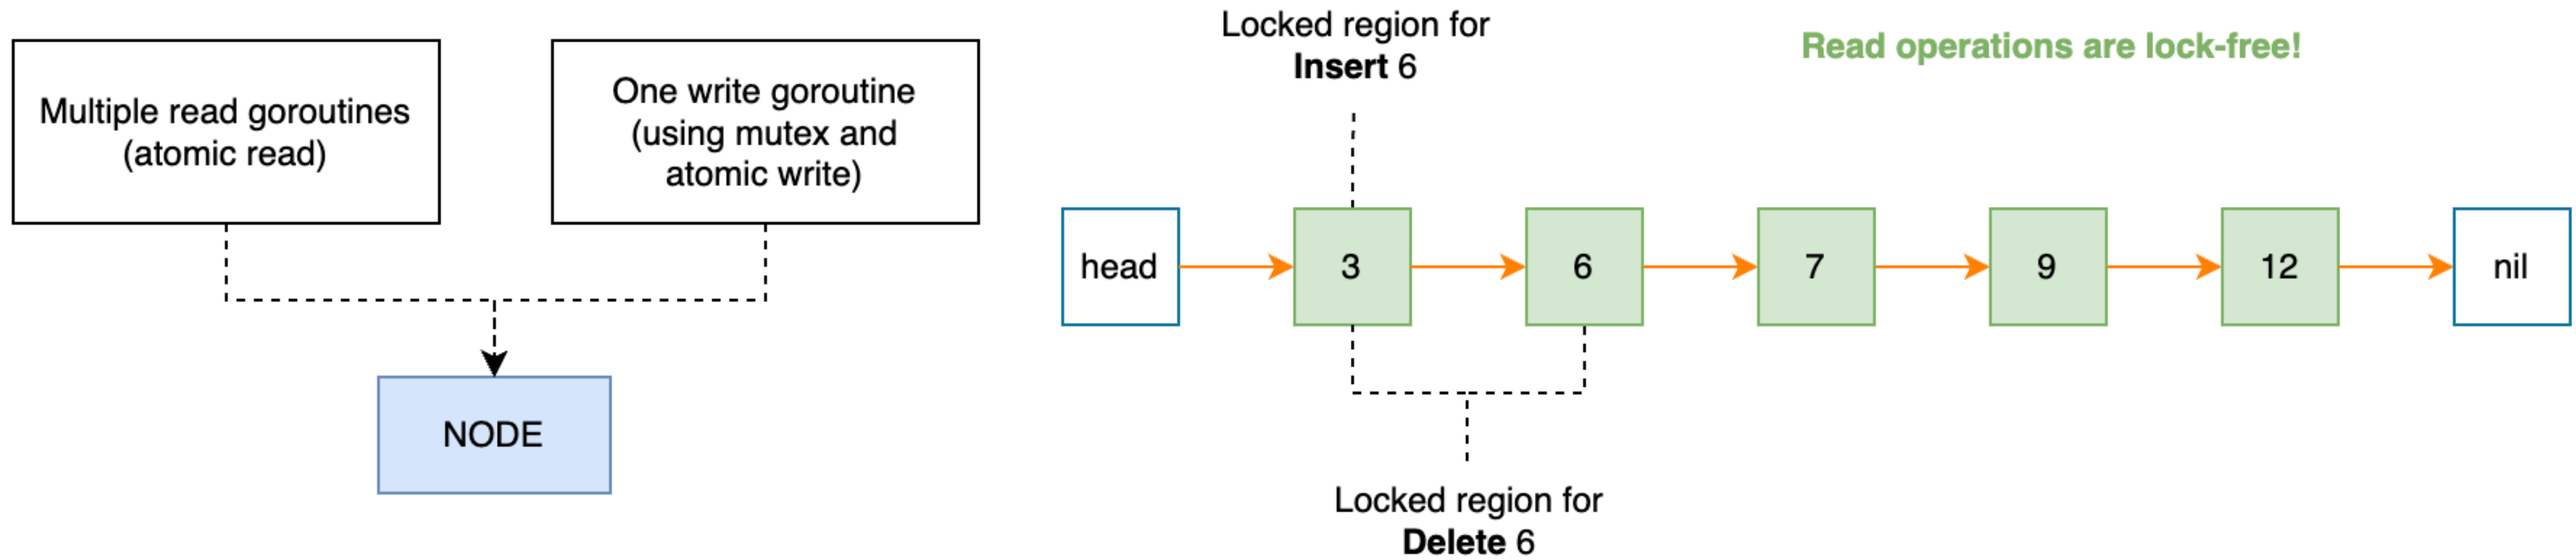
We upgraded a single read-write mutex map to a shared map (multiple read-write mutex maps), and by **reducing the locked regions for each write operation**, we achieved a 2x performance improvement.

**Let's try one item per sub-map**





Sub-map (or node, with only one item). Each time a new key-value pair is inserted, we create a new node, and all nodes are connected through pointers, which is actually a **linked list**.

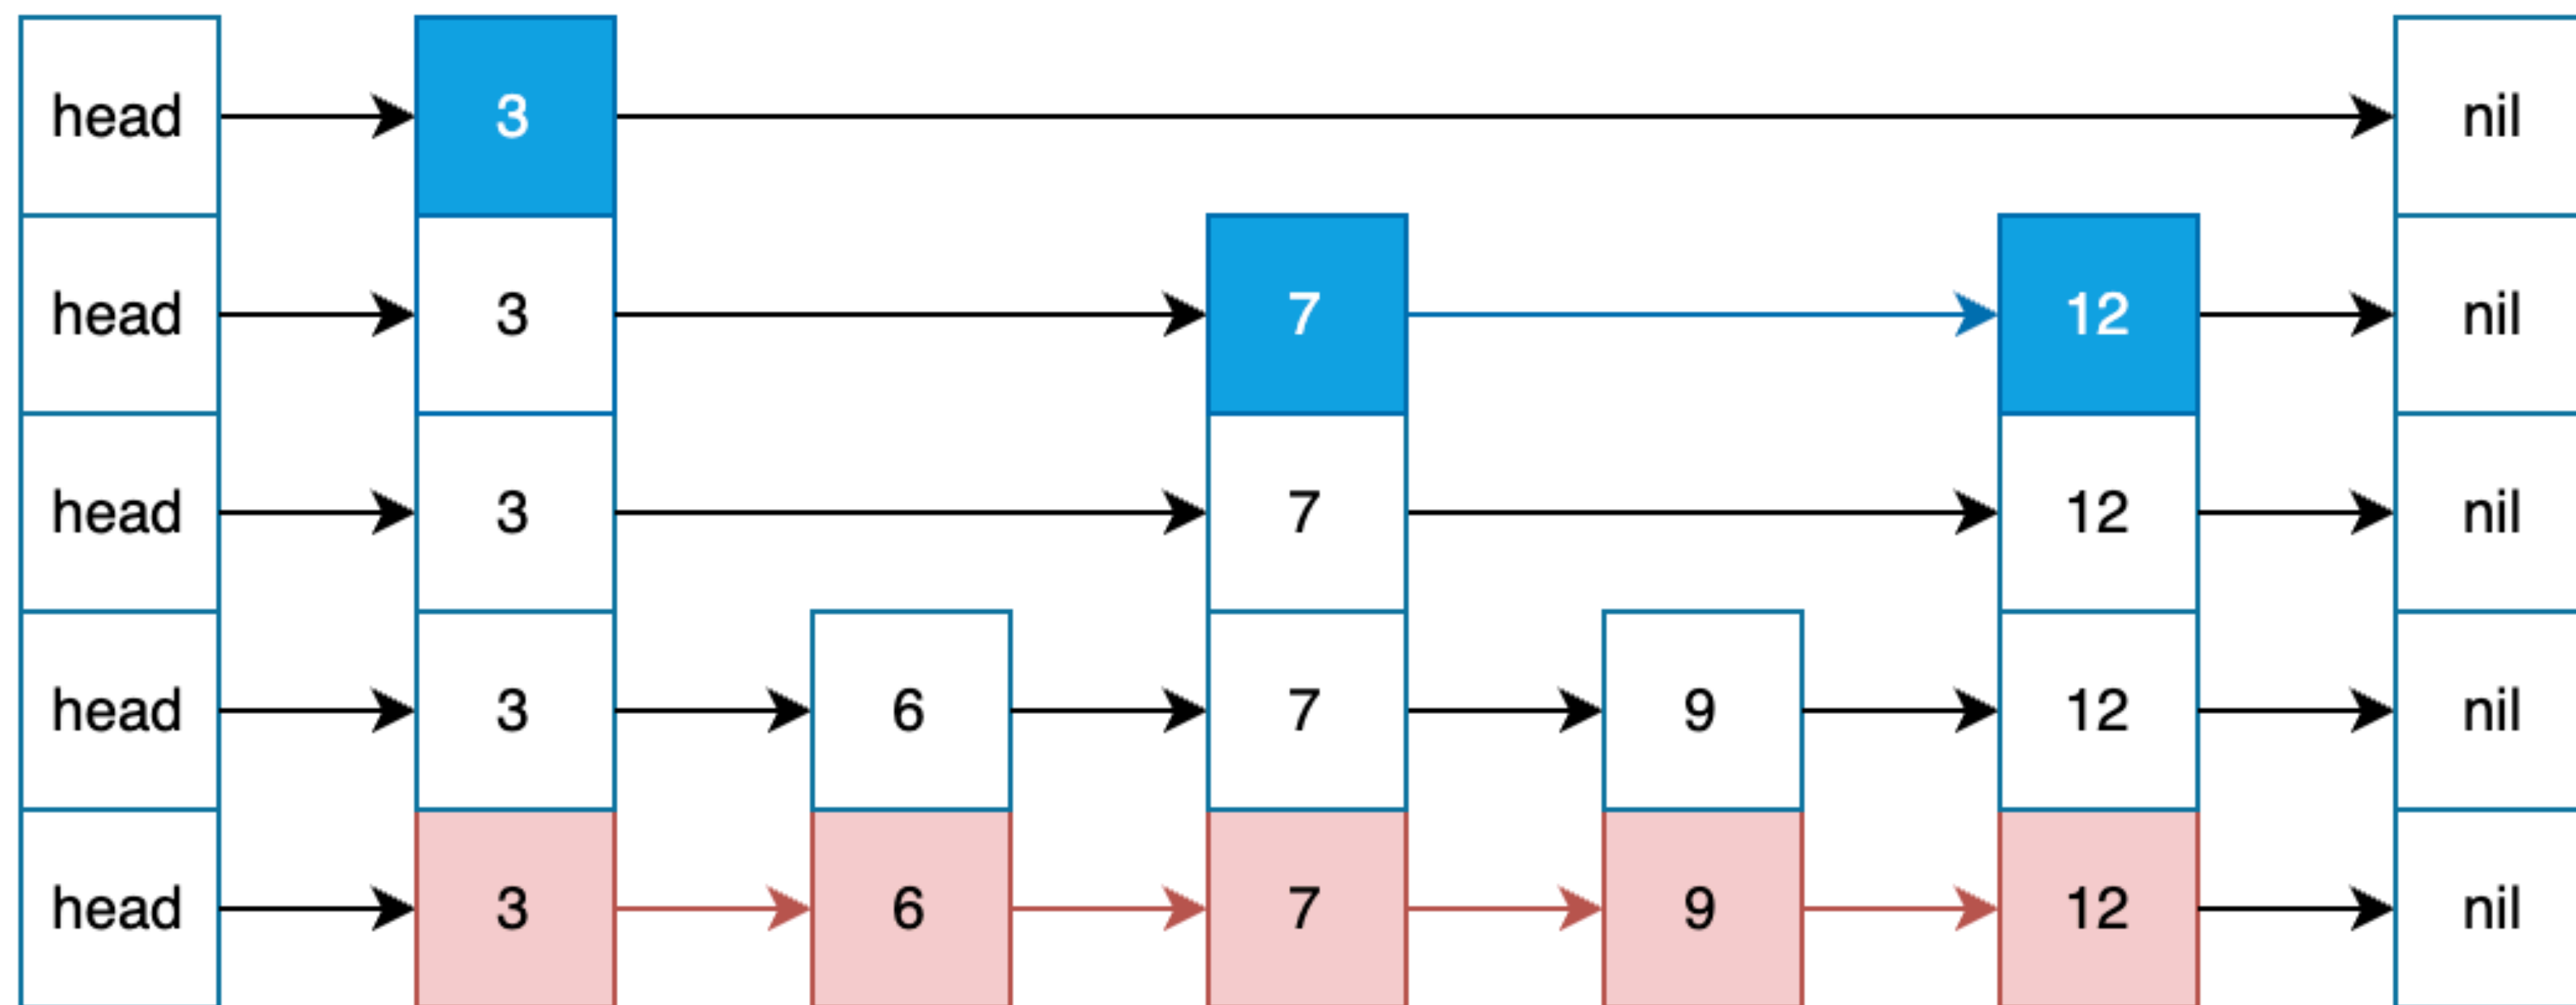


- \*All write operations must acquire the lock
- \*All read operations are lock-free



**Ordered concurrent-safe linked list**

Linked list accessing time complexity:  $O(n)$

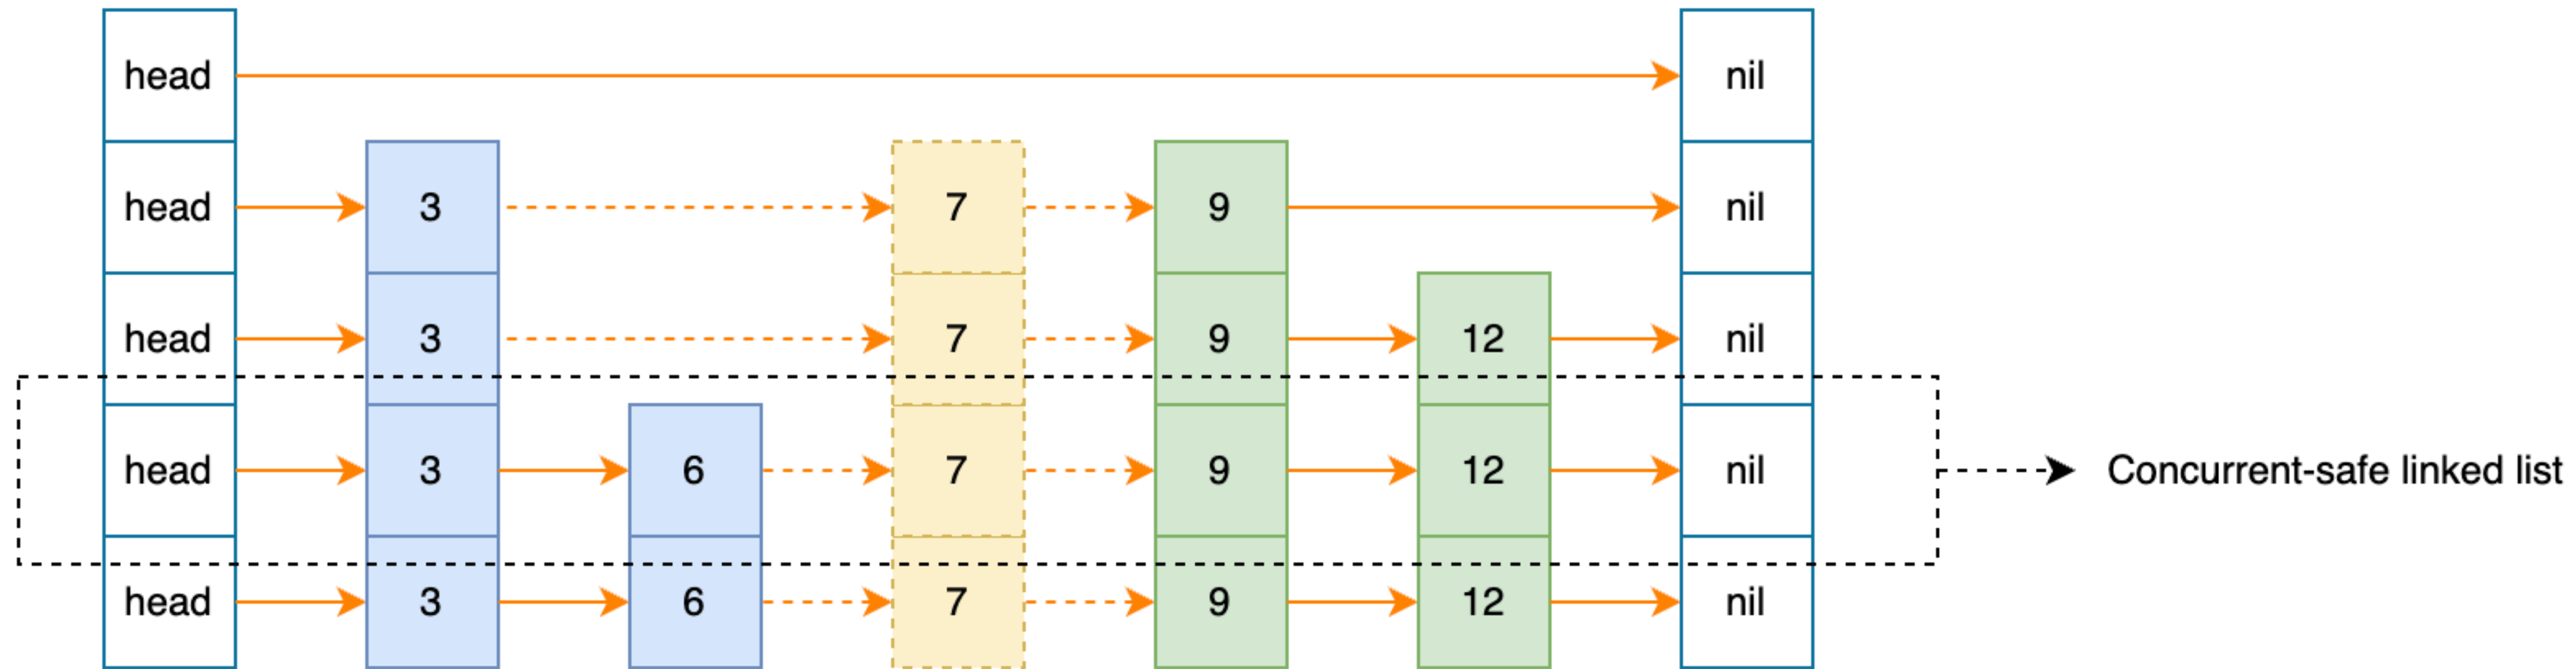


Linked List Search Process  $O(n)$



Skip List Search Process  $O(\log n)$

SkipList = multiple linked lists



Skipmap = multiple concurrent-safe linked lists

	syncmap.txt	skipmap.txt
	sec/op	sec/op vs base
LoadSize1000-16	7.151n ± 23%	5.527n ± 15% -22.71% (p=0.002 n=10)
Store-16	913.4n ± 3%	142.2n ± 5% -84.44% (p=0.000 n=10)
70Load30Store-16	699.55n ± 5%	62.84n ± 6% -91.02% (p=0.000 n=10)
90Load9Store1Delete-16	595.50n ± 1%	32.10n ± 5% -94.61% (p=0.000 n=10)
90Load8Store1Delete1Range-16	40594.0n ± 1%	397.3n ± 2% -99.02% (p=0.000 n=10)
geomean	643.6n	57.52n -91.06%

skipmap is ~20x faster in typical case(90%Load9%Store1%Delete)

	Time complexity (read and write)
skipmap	$O(\log n)$
sync.Map	$O(1)$

skipmap is faster in high-concurrency read-write cases

sync.Map may be faster in low-concurrency cases

## Reduce the locked regions for each concurrent operation

- \* Read-write mutex map -> entire map
- \* Shared map -> one sub-map
- \* Skipmap -> few items